

Recap

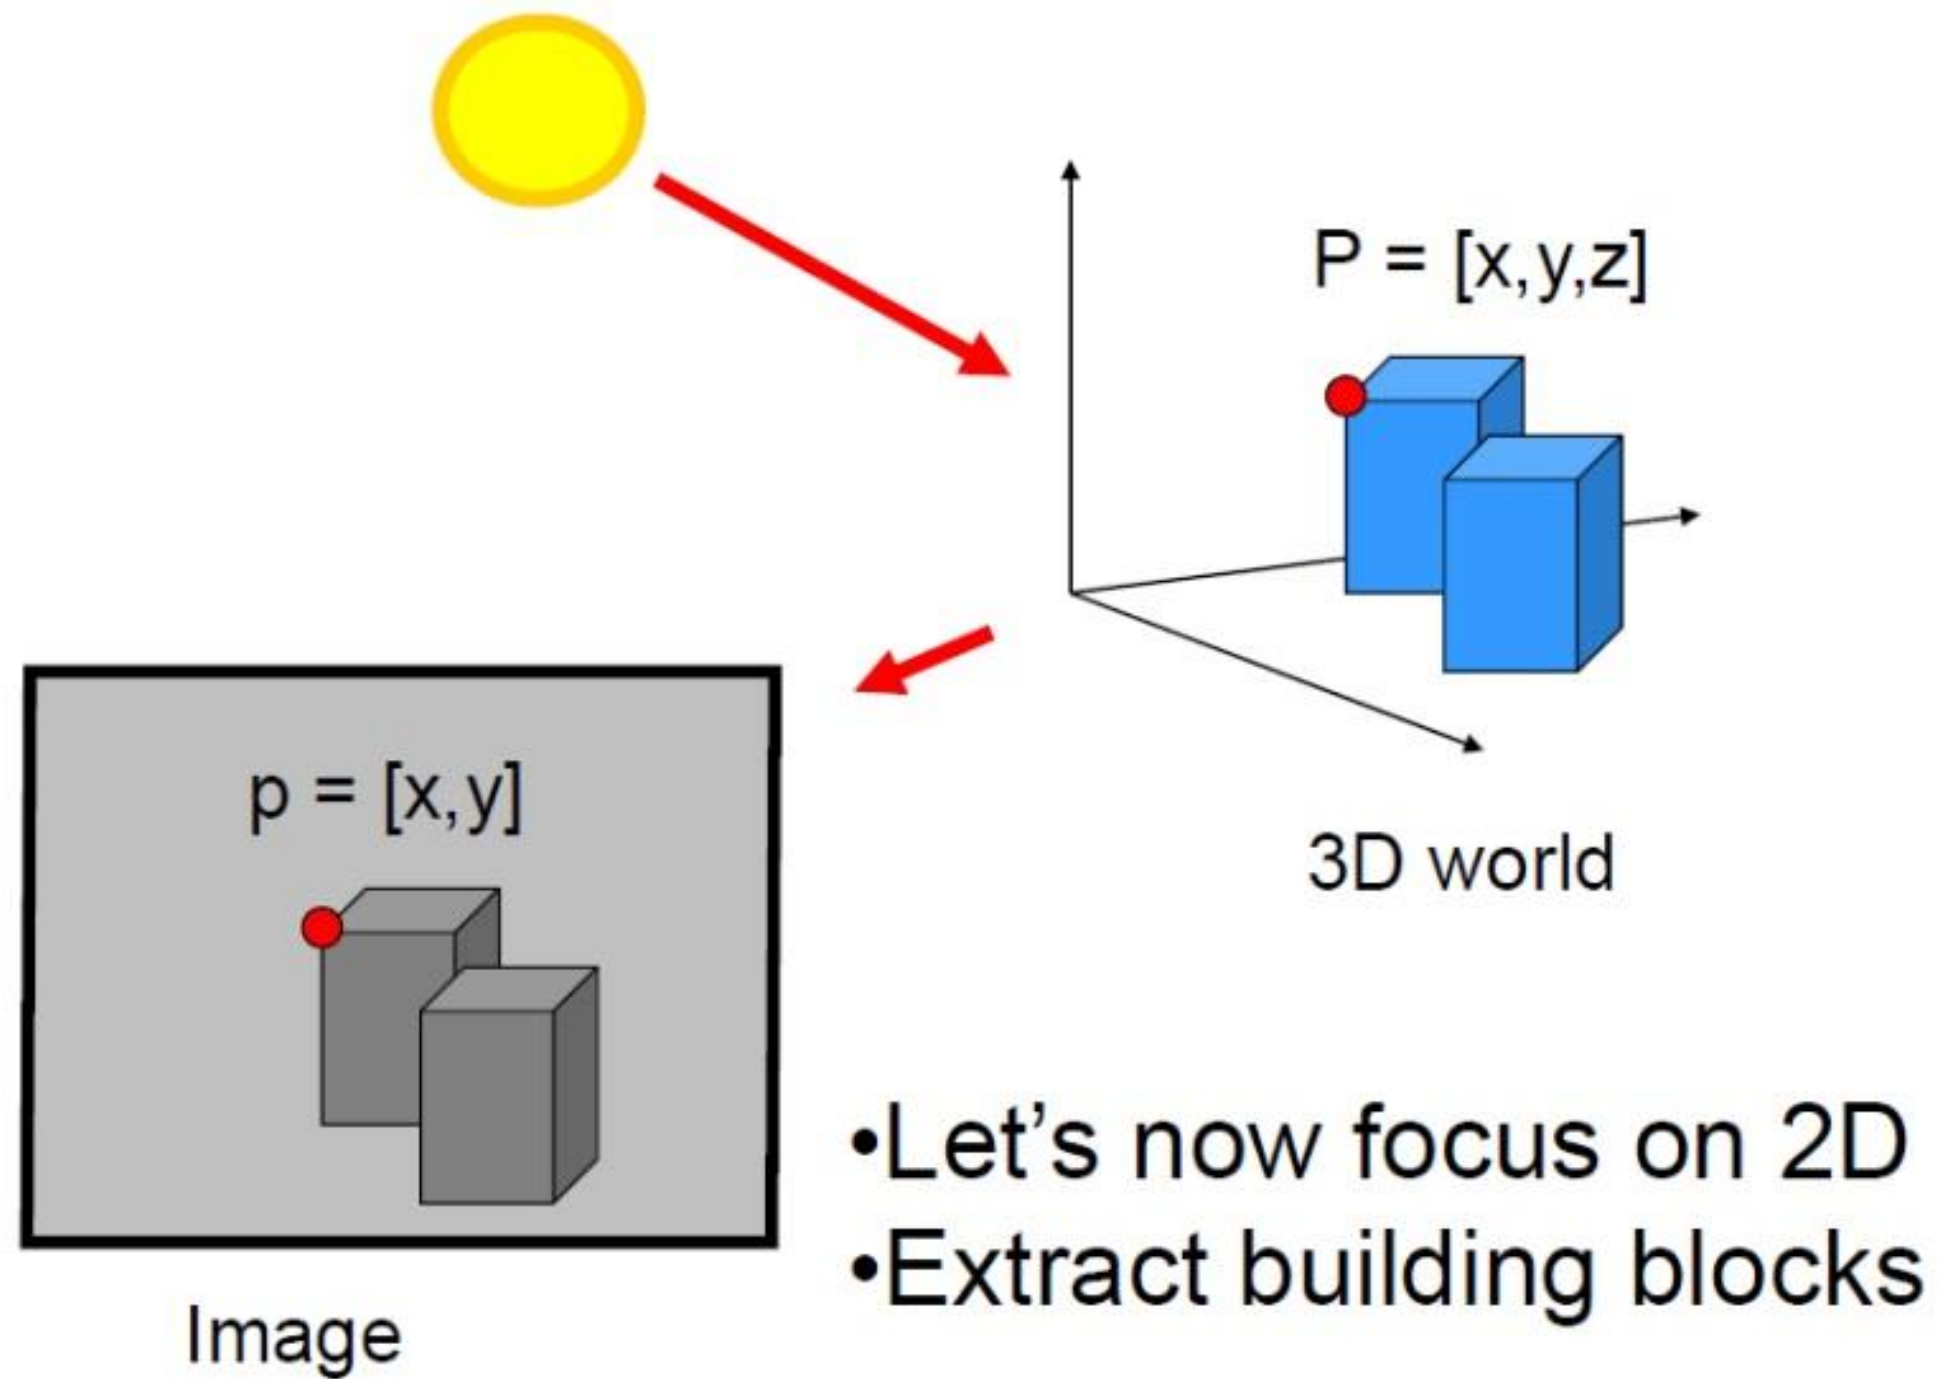


Image filtering

- Image filtering: compute function of local neighborhood at each position
- Really important!
 - Enhance images
 - Denoise, resize, increase contrast, etc.
 - Extract information from images
 - Texture, edges, distinctive points, etc.
 - Detect patterns
 - Template matching

Image filtering - Example

$g[\cdot, \cdot]$

1	1	1
1	1	1
1	1	1

Image filtering - Example

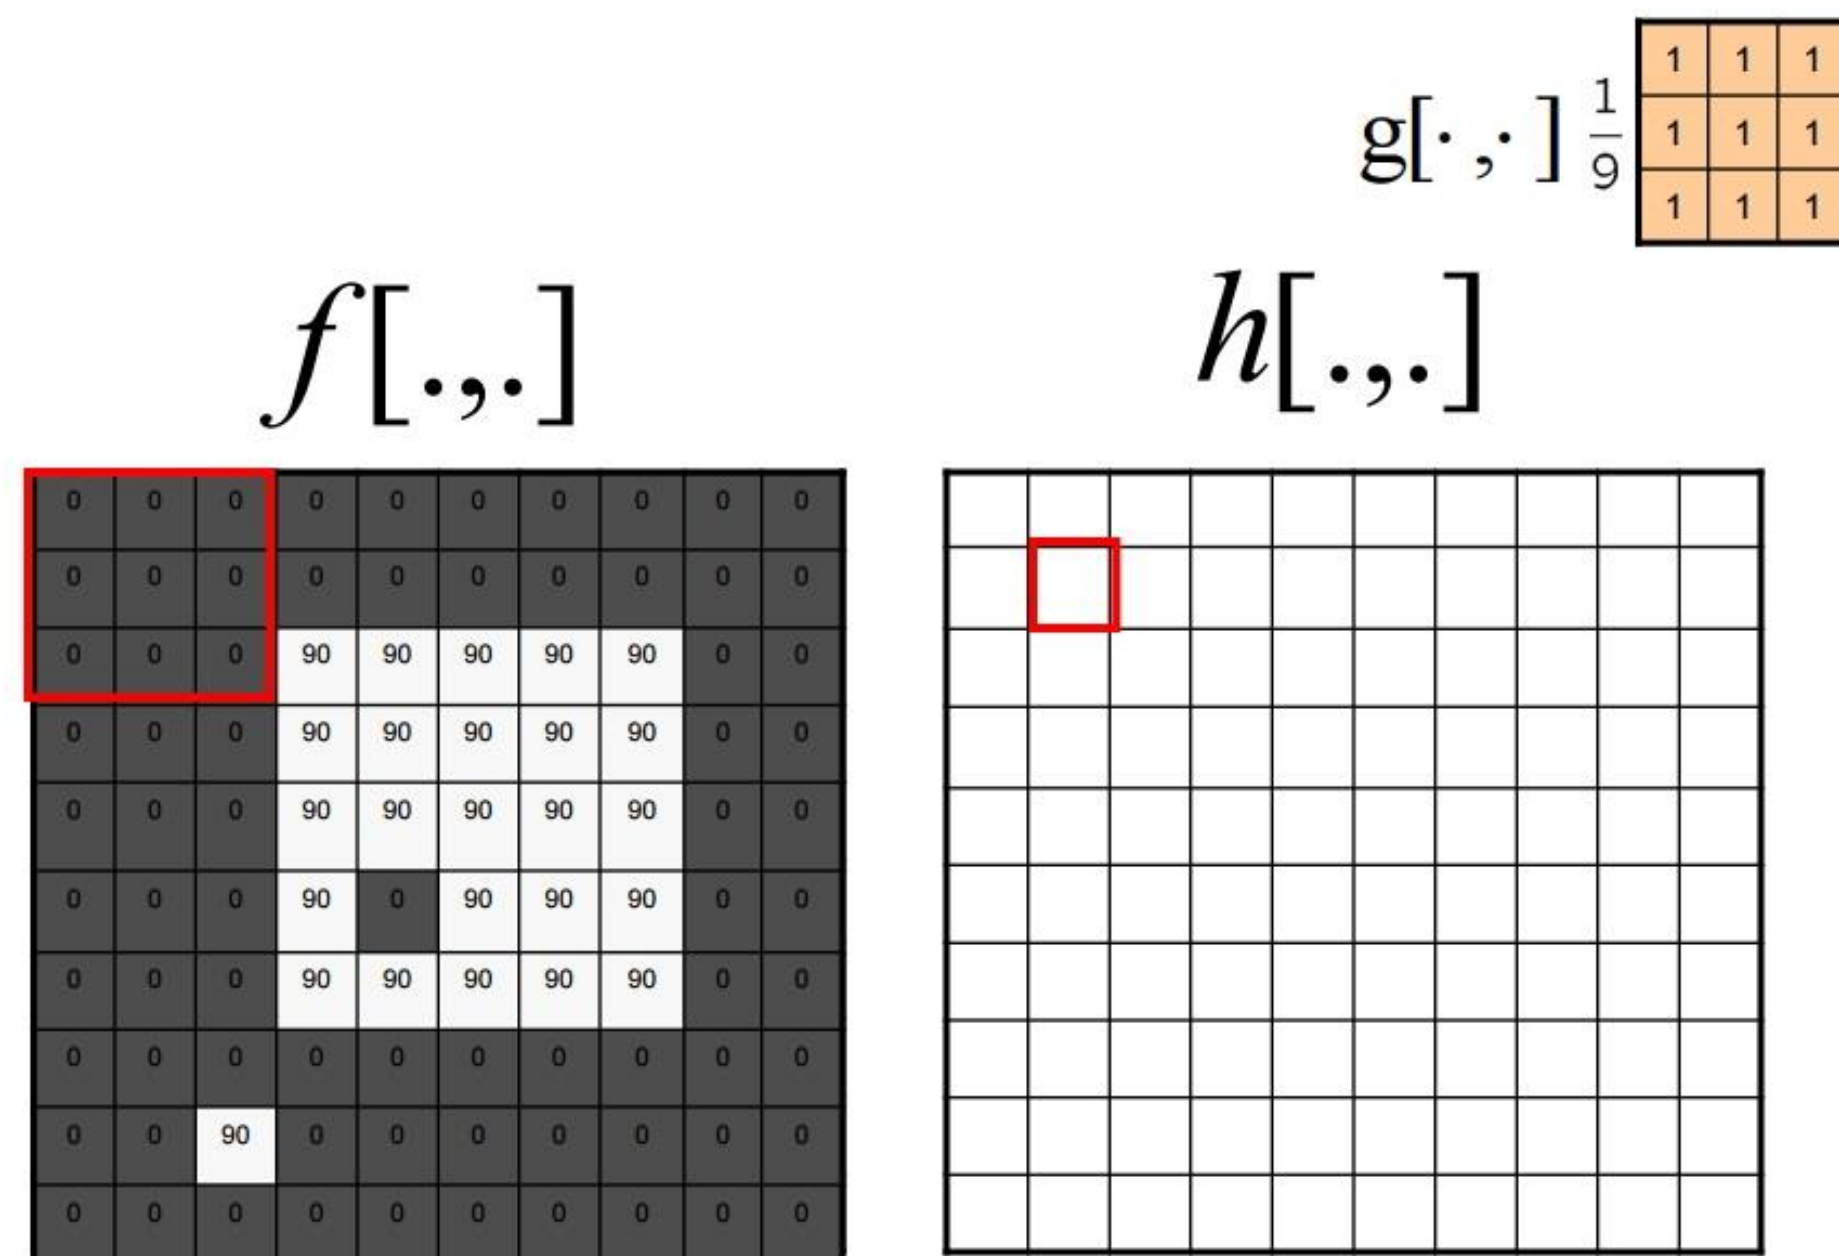


Image filtering - Example

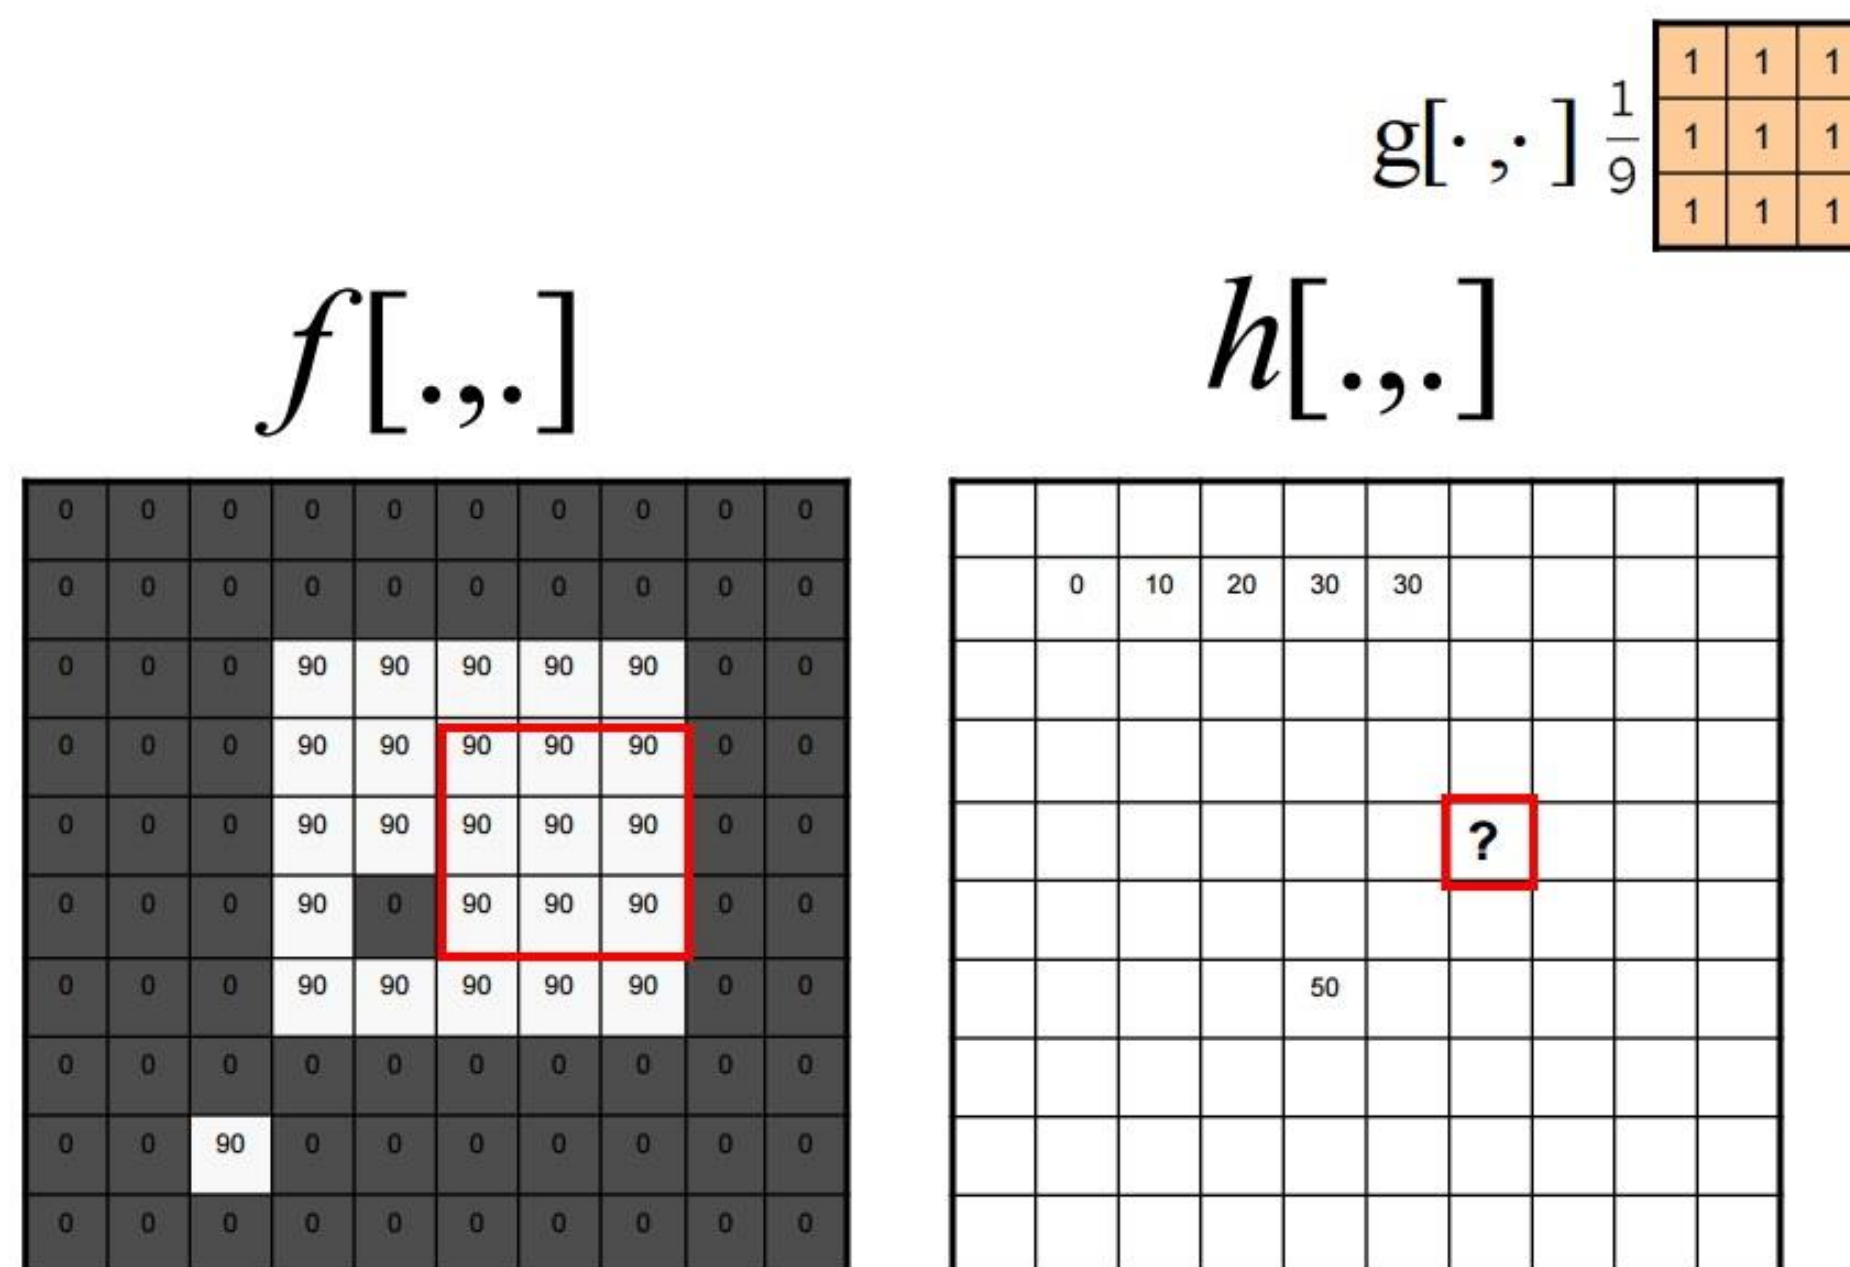


Image filtering – Correlation Formula

$$g[\cdot, \cdot] \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

$$f[\cdot, \cdot]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$h[\cdot, \cdot]$$

	0	10	20	30	30	30	20	10	
	0	20	40	60	60	60	40	20	
	0	30	60	90	90	90	60	30	
	0	30	50	80	80	90	60	30	
	0	30	50	80	80	90	60	30	
	0	20	30	50	50	60	40	20	
	10	20	30	30	30	30	20	10	
	10	10	10	0	0	0	0	0	

$$h[m, n] = \sum_{k, l} g[k, l] f[m + k, n + l]$$

Image filtering - Example

What does it do?

- Replaces each pixel with an average of its neighborhood
- Achieve smoothing effect (remove sharp features)

$$\frac{1}{9} g[\cdot, \cdot]$$

1	1	1
1	1	1
1	1	1

Image filtering - Example

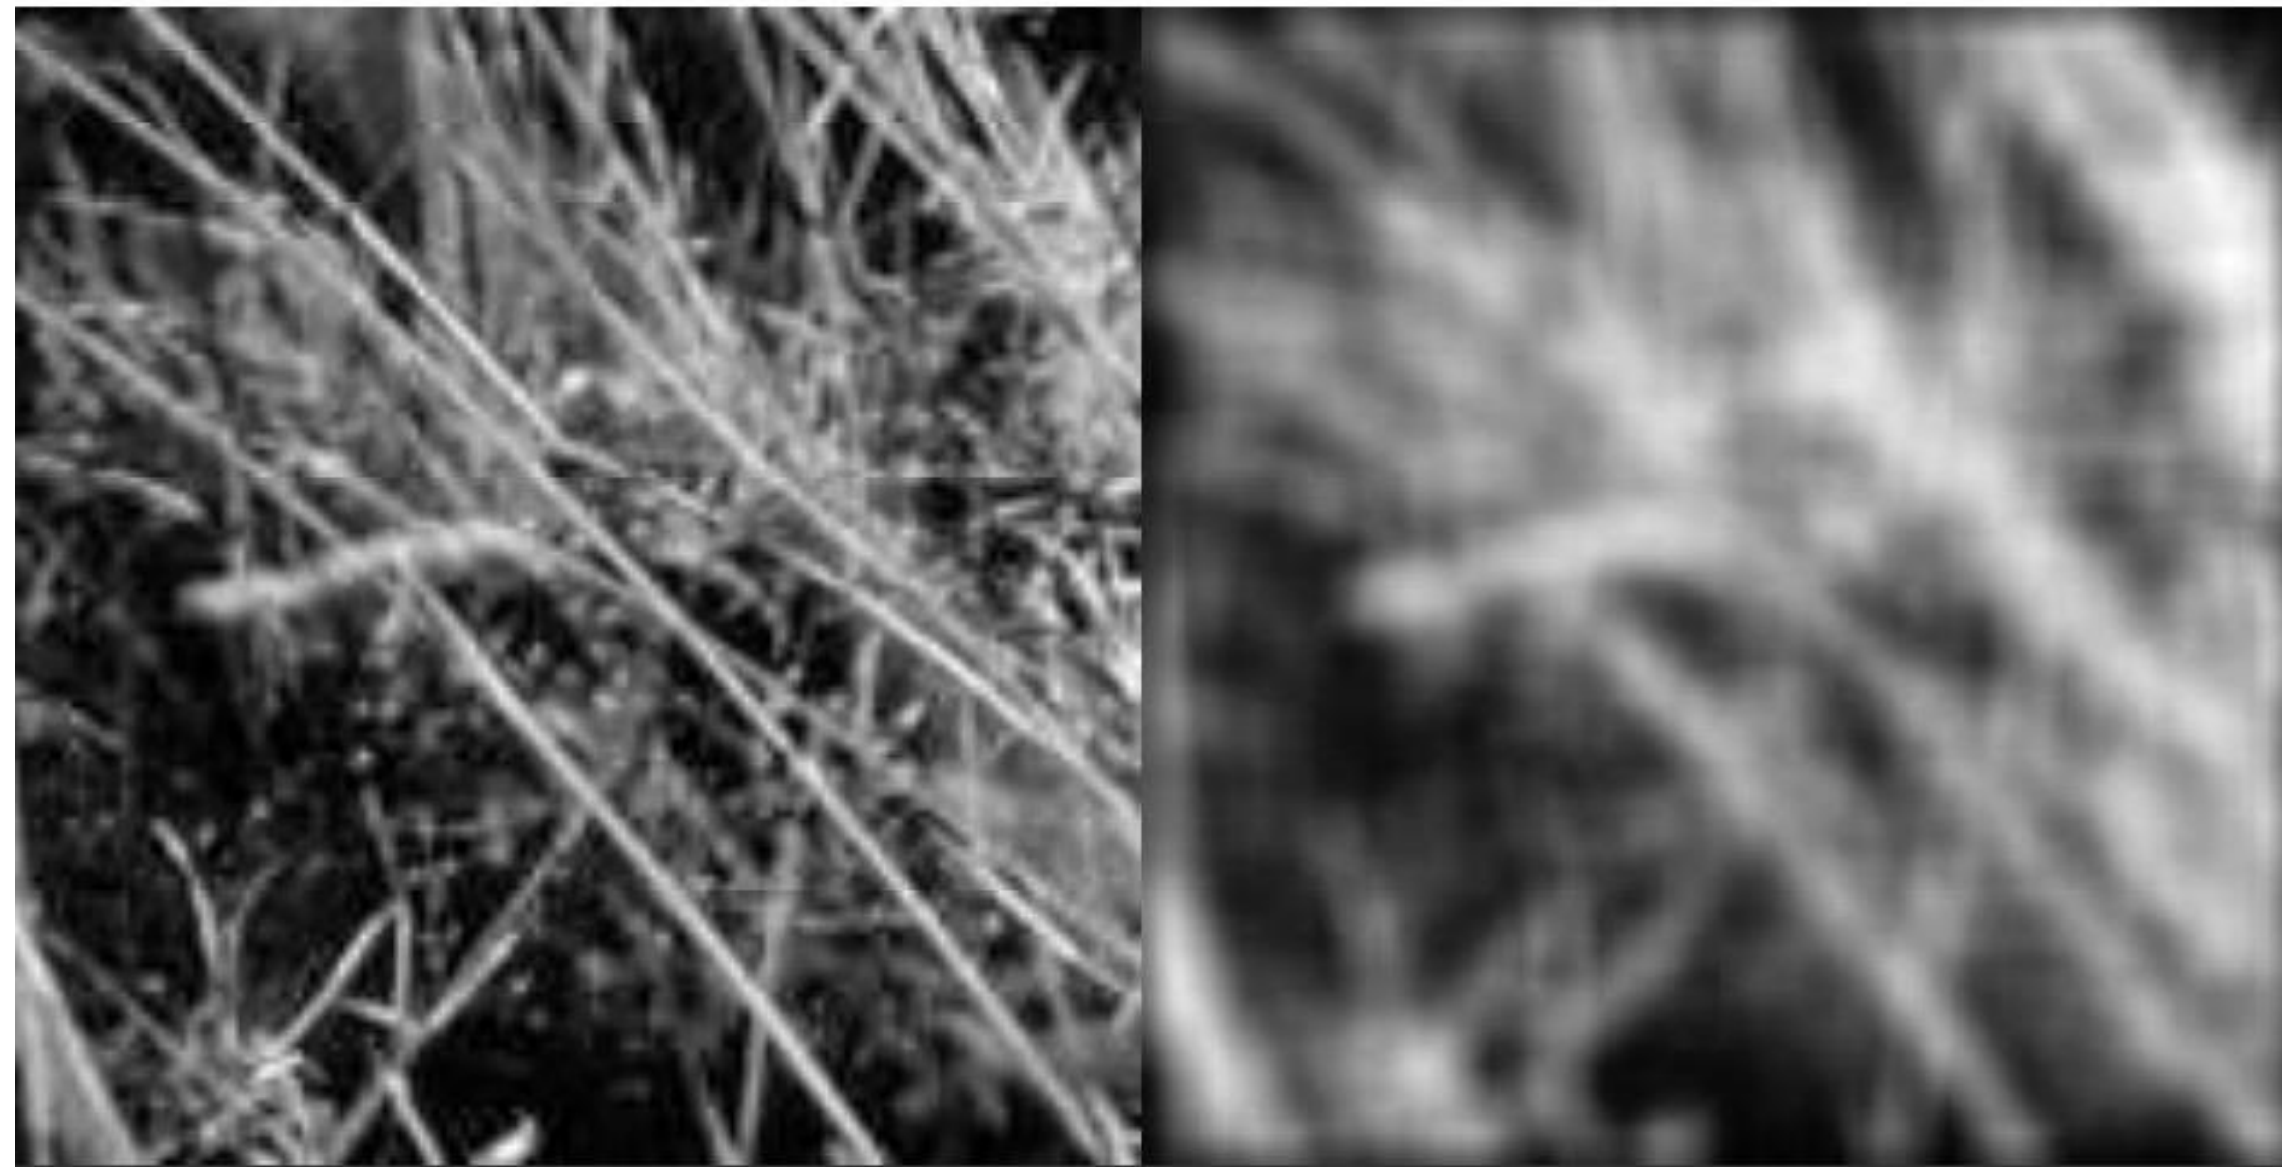


Image filtering – Using Correlation Formula



Original

0	0	0
0	1	0
0	0	0

?

Image filtering - Questions



Original

0	0	0
0	1	0
0	0	0



Filtered
(no change)

Image filtering - Questions



Original

0	0	0
0	0	1
0	0	0

?

Source: D. Lowe

Image filtering - Questions

Correlation formula leads to counter intuitive image filtering



original

0	0	0
0	0	1
0	0	0



shifted

Image filtering - Convolution

- Same as cross-correlation, except that the kernel is “flipped” (horizontally and vertically)

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i - u, j - v]$$

This is called a **convolution** operation:

$$G = H * F$$

- Convolution is **commutative** and **associative**

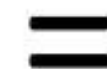
Image filtering - Convolution



Original



0	0	0
1	0	0
0	0	0



Shifted left
By 1 pixel

Image filtering - Sharpening



Original

0	0	0
0	2	0
0	0	0

—

$\frac{1}{9}$

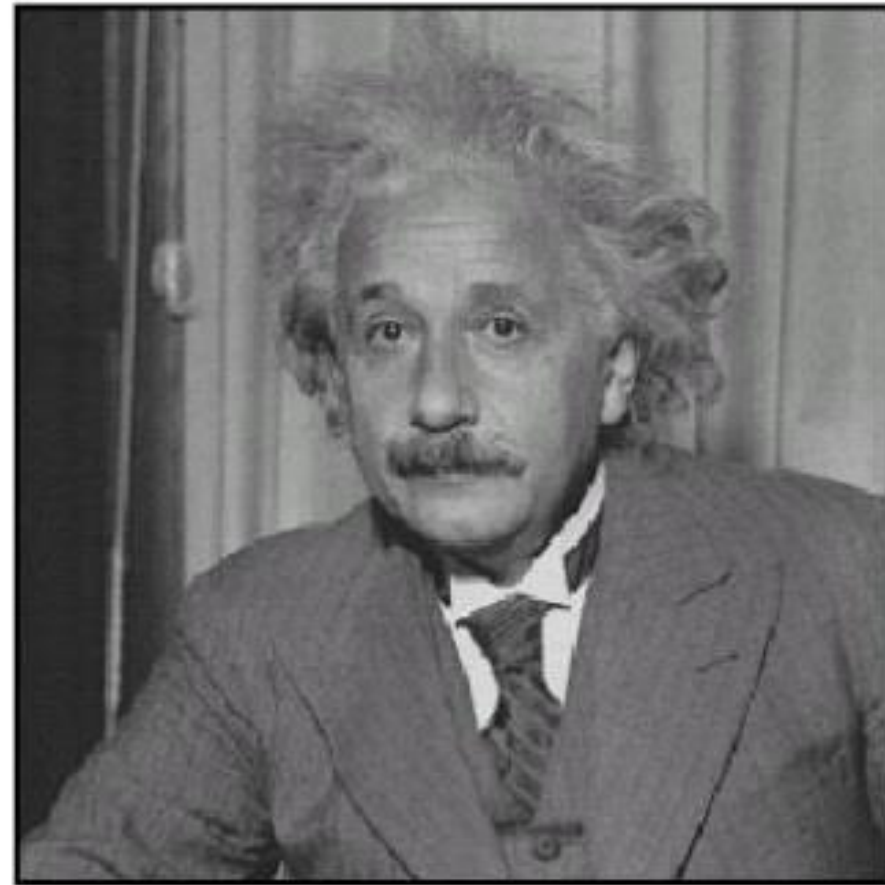
1	1	1
1	1	1
1	1	1



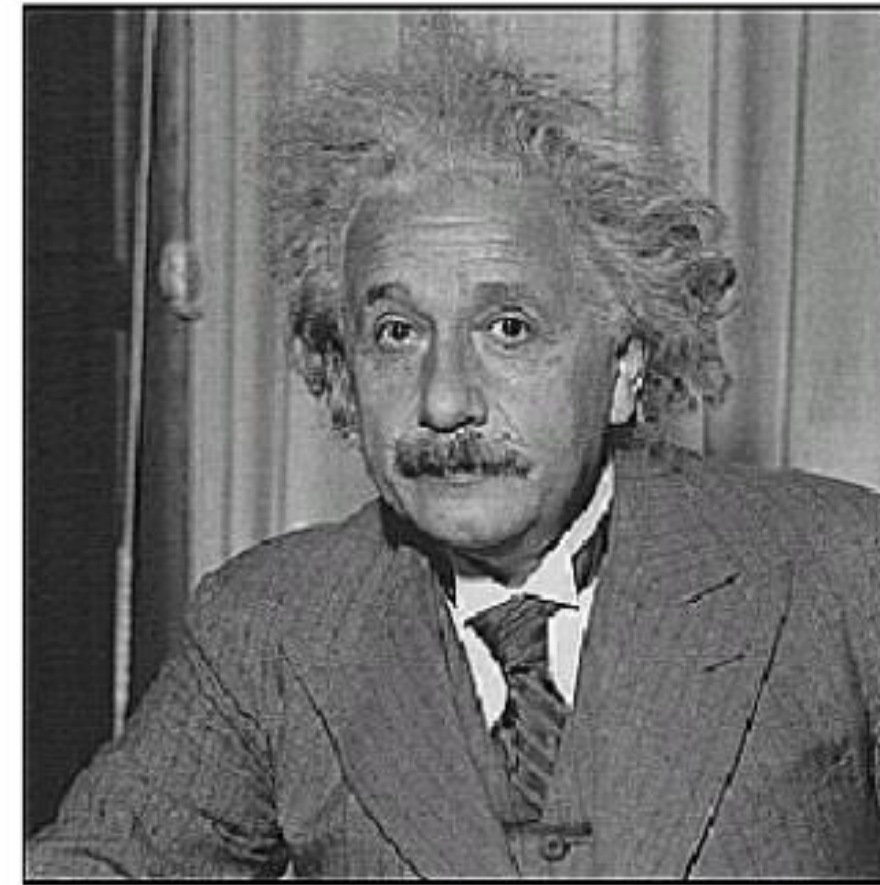
Sharpening filter

- Accentuates differences with local average

Image filtering - Sharpening

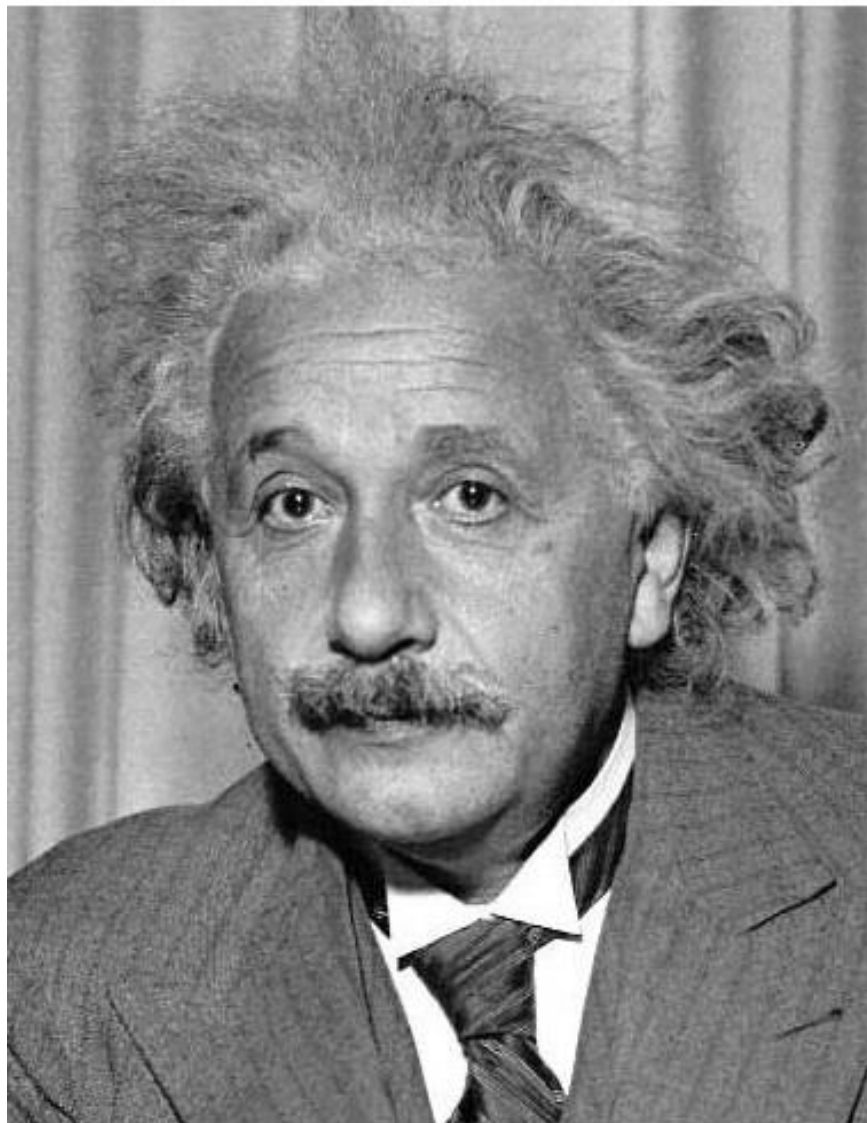


before



after

Image filtering - Edges



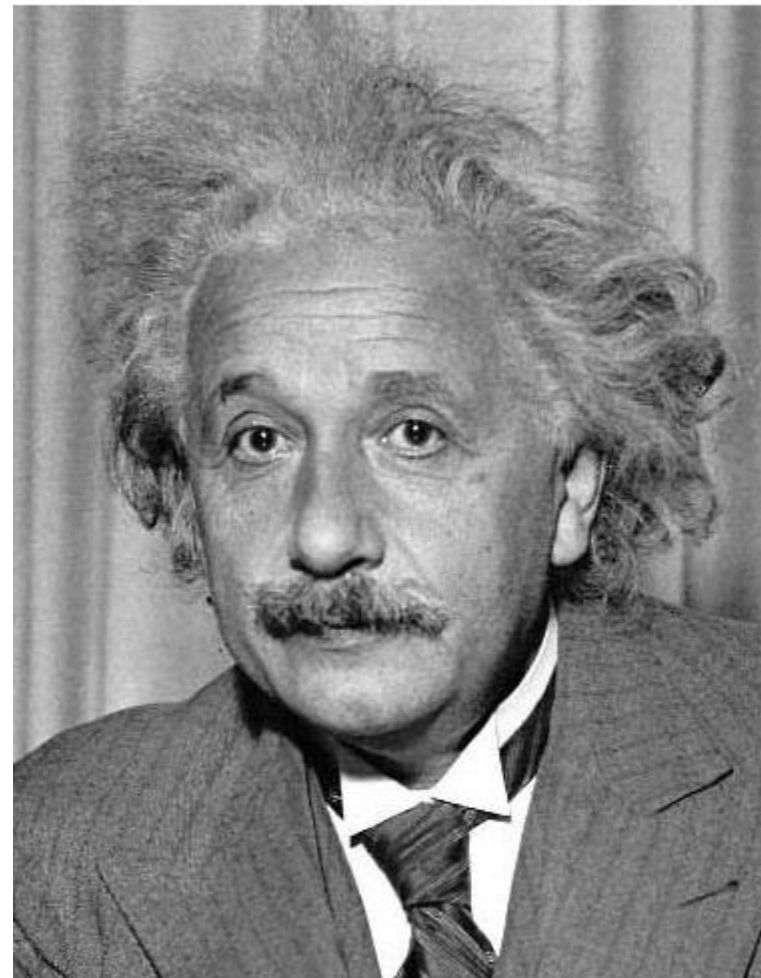
1	0	-1
2	0	-2
1	0	-1

Sobel



Vertical Edge
(absolute value)

Image filtering - Edges



1	2	1
0	0	0
-1	-2	-1

Sobel



Horizontal Edge
(absolute value)

Image filters - Separability

Are you familiar with Matrix outer product?

2D convolution
(center location only)

$$\begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} * \begin{bmatrix} 2 & 3 & 3 \\ 3 & 5 & 5 \\ 4 & 4 & 6 \end{bmatrix}$$

The filter factors
into a product of 1D
filters:

$$\begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} \times \begin{bmatrix} 1 & 2 & 1 \end{bmatrix}$$

Perform convolution
along rows:

$$\begin{bmatrix} 1 & 2 & 1 \end{bmatrix} * \begin{bmatrix} 2 & 3 & 3 \\ 3 & 5 & 5 \\ 4 & 4 & 6 \end{bmatrix} = \begin{bmatrix} & 11 & \\ & 18 & \\ & 18 & \end{bmatrix}$$

Followed by convolution
along the remaining column:

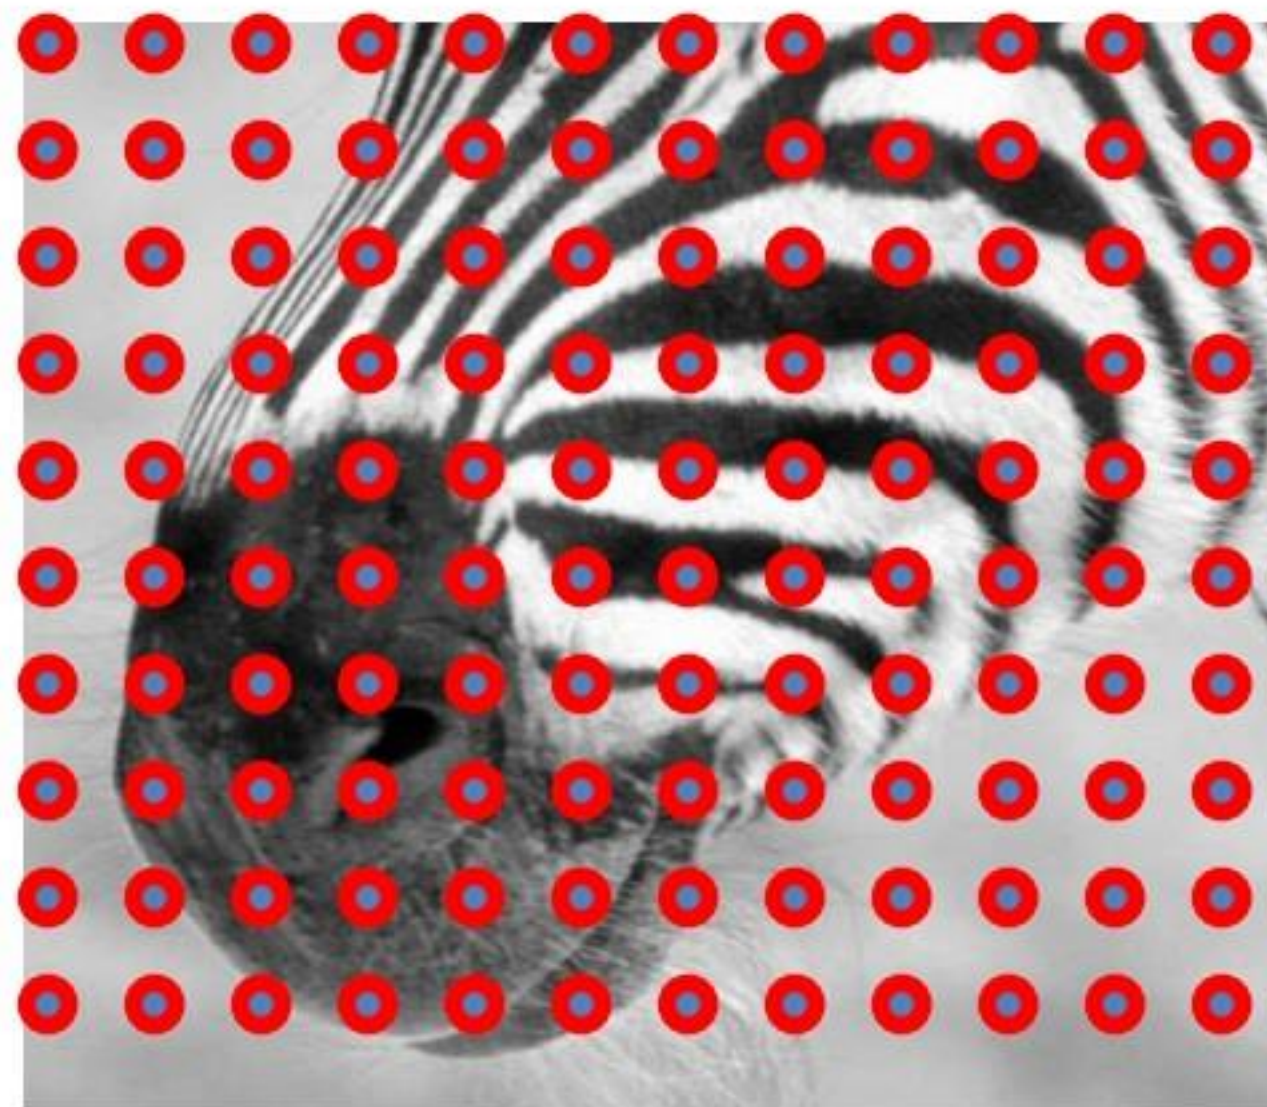
Image Resizing

Why does a lower resolution image still make sense to us? What do we lose?



Image Resizing – Basic Approach

Do you see any potential Issues?



Throw away every other row and
column to create a 1/2 size image

Image Resizing – Aliasing

- 1D example (sinewave):

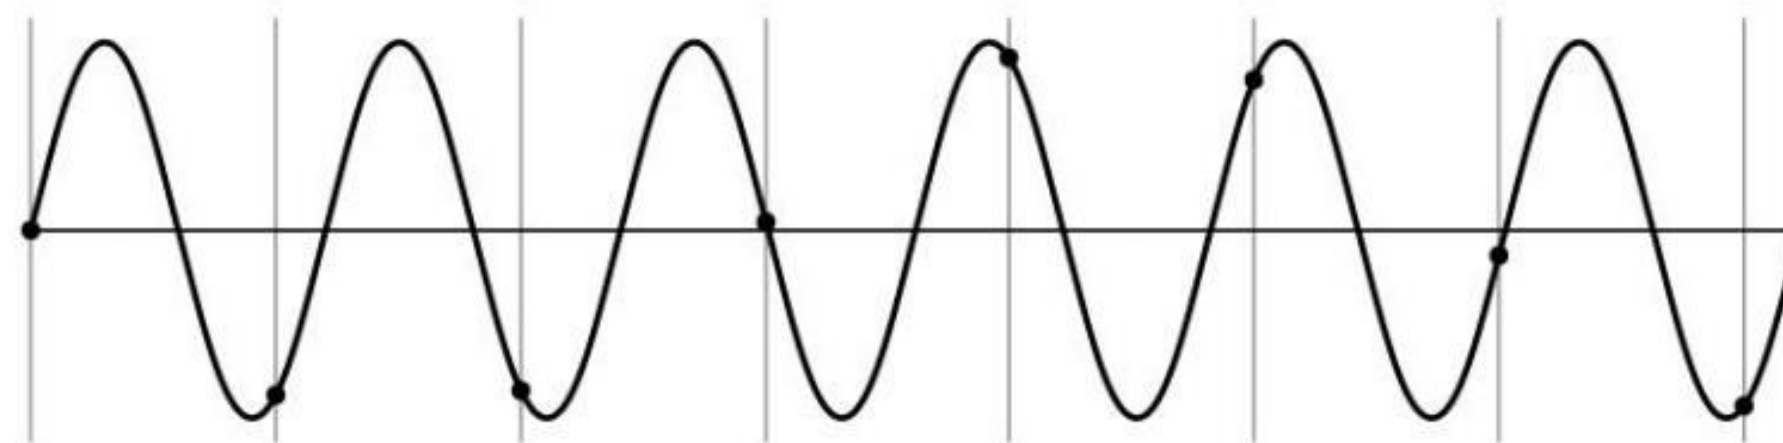


Image Resizing – Aliasing

Reconstructed signal not the same as original!

- 1D example (sinewave):

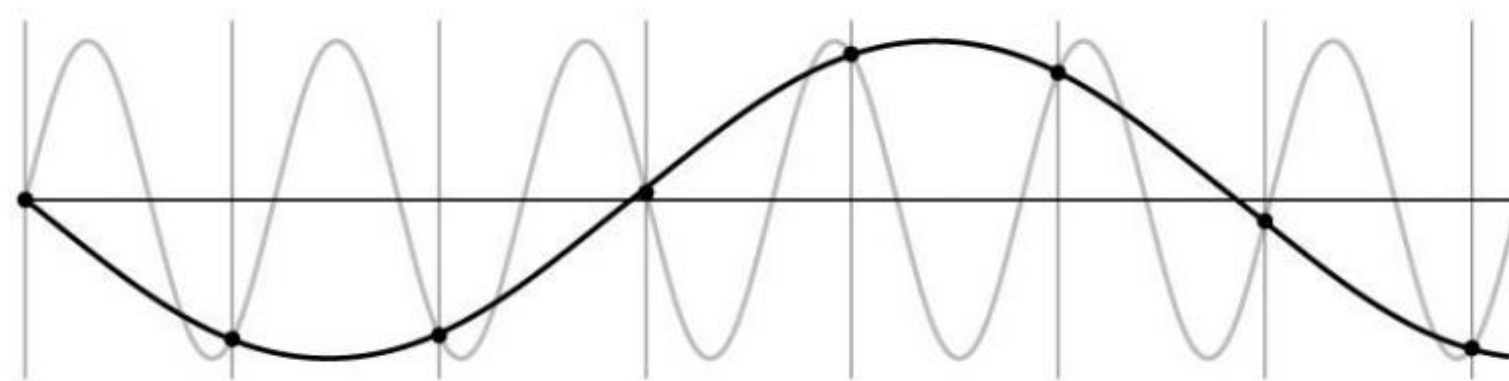
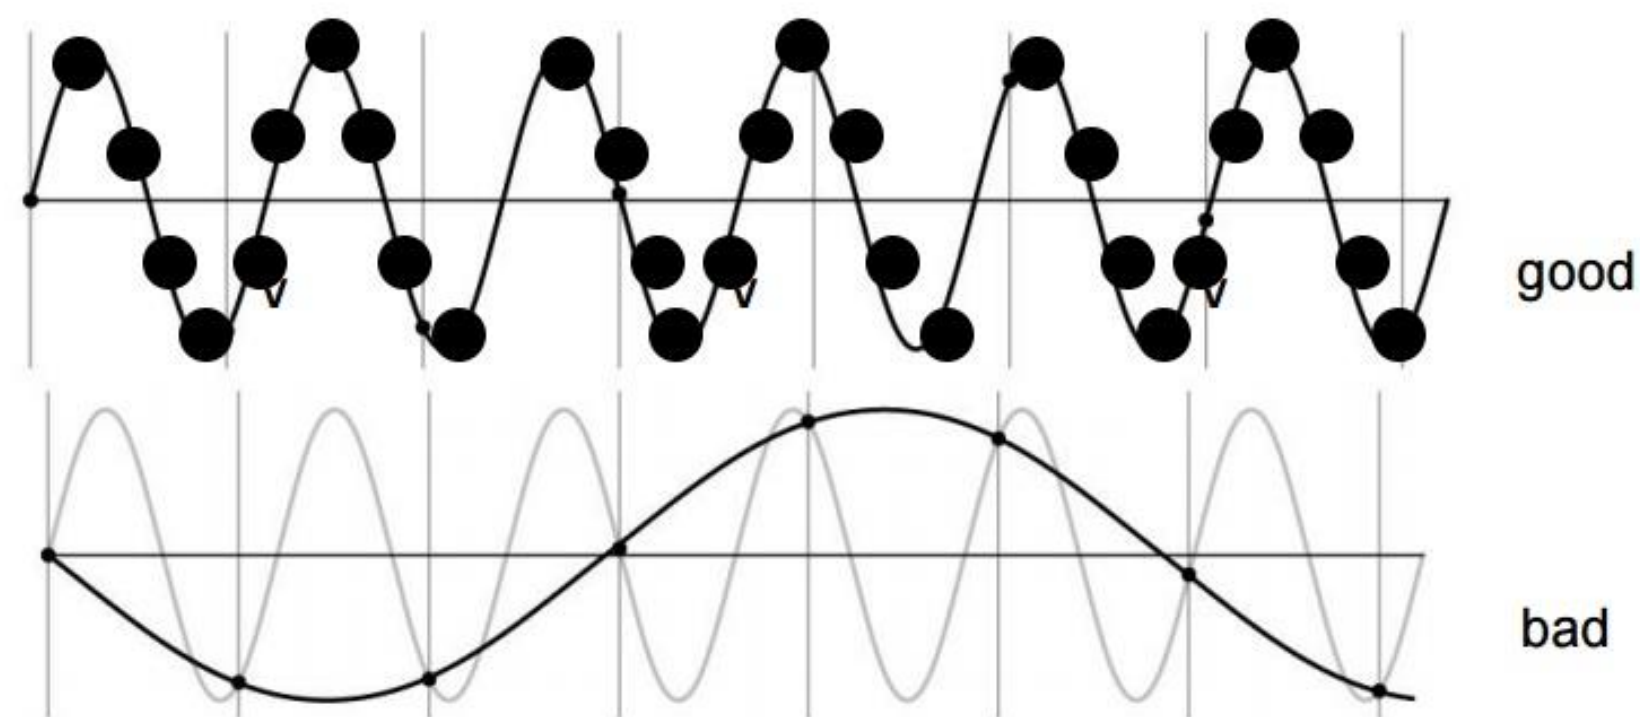
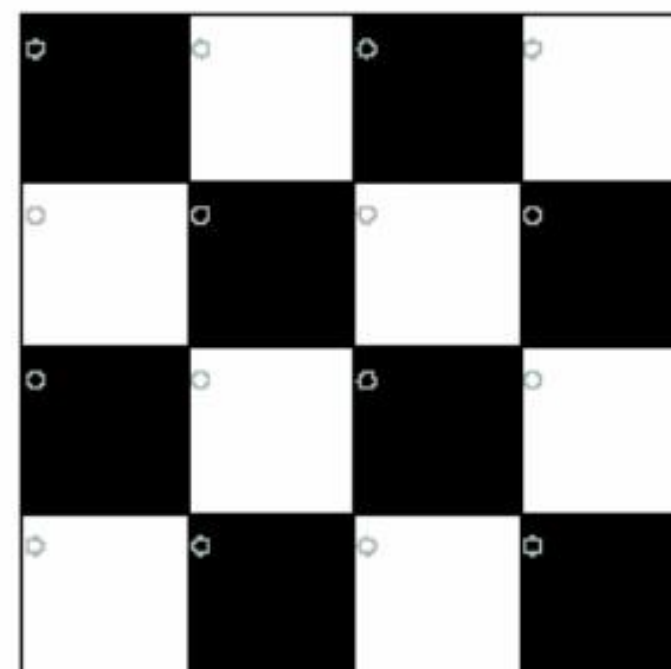
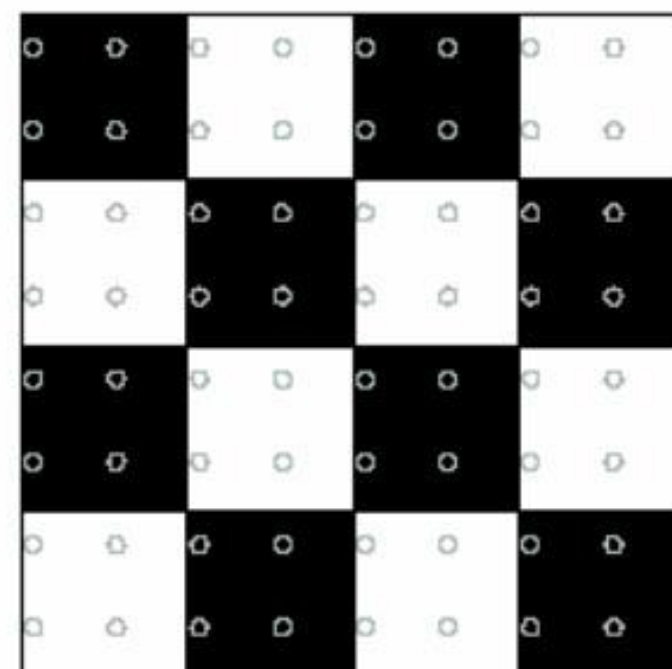


Image Resizing – Aliasing

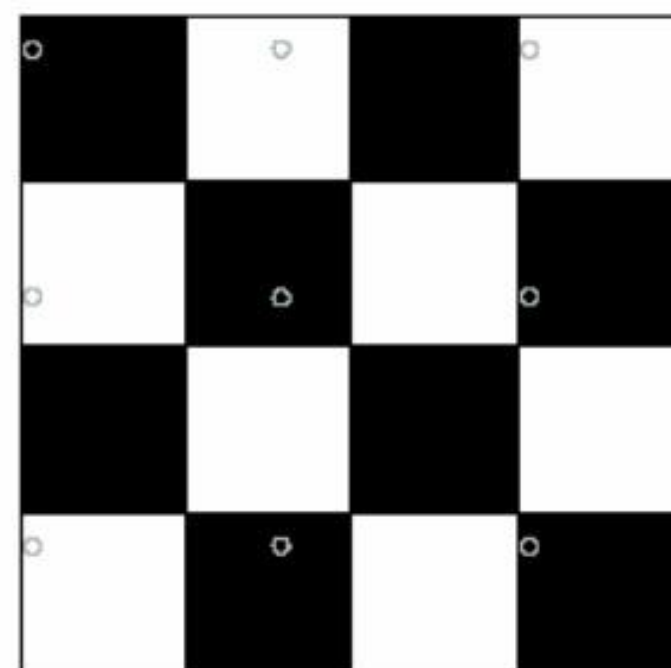
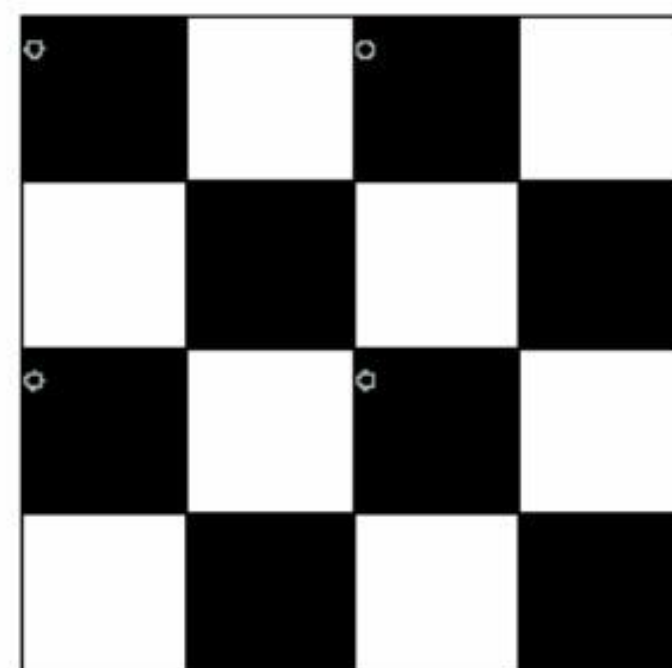
- When sampling a signal at discrete intervals, the sampling frequency must be $\geq 2 \times f_{\max}$
- f_{\max} = max frequency of the input signal
- This will allow to reconstruct the original perfectly from the sampled version



Sampling



Good sampling:
• Sample often or,
• Sample wisely



Bad sampling:
• see aliasing in action!

How to Subsample Images

Solutions:

- Sample more often
- Get rid of all frequencies that are greater than half the new sampling frequency
 - Will lose information
 - But it's better than aliasing
 - Apply a smoothing filter

Image Sub-Sampling Pipeline

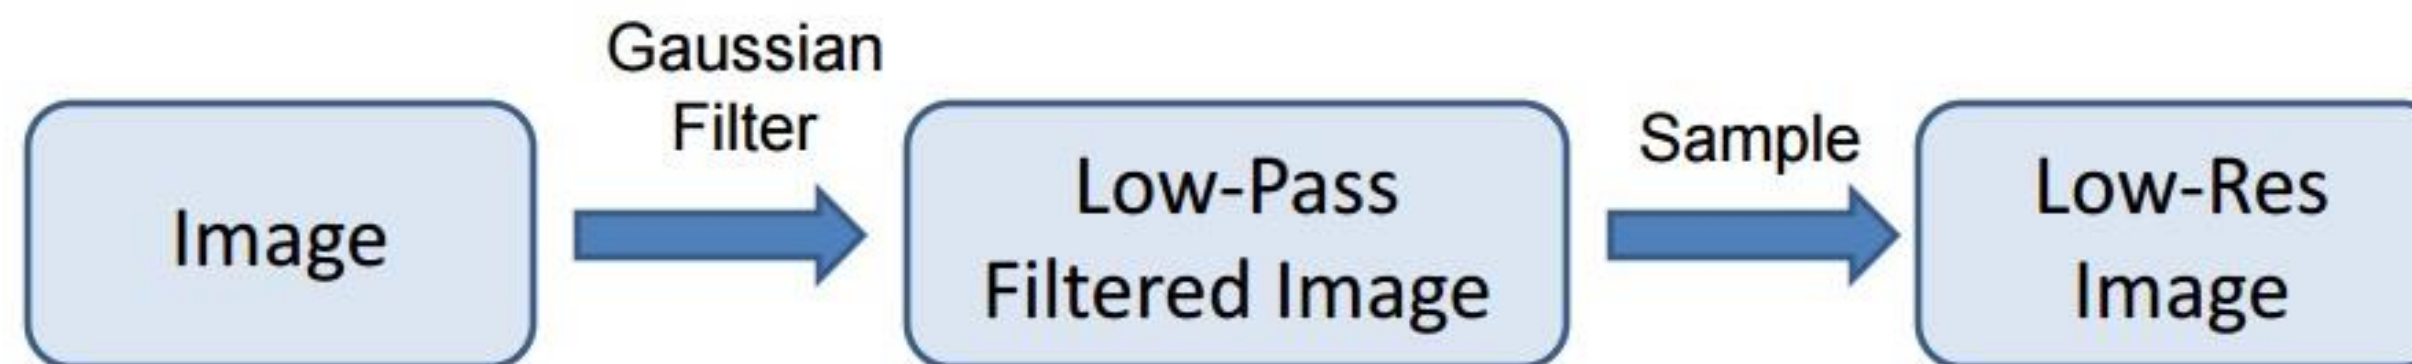


Image Fourier Analysis

Any signal can be approximate by a summation of sinusoids!

Our building block:

$$A \sin(\omega x + \phi)$$

Add enough of them to get any signal $g(x)$ you want!

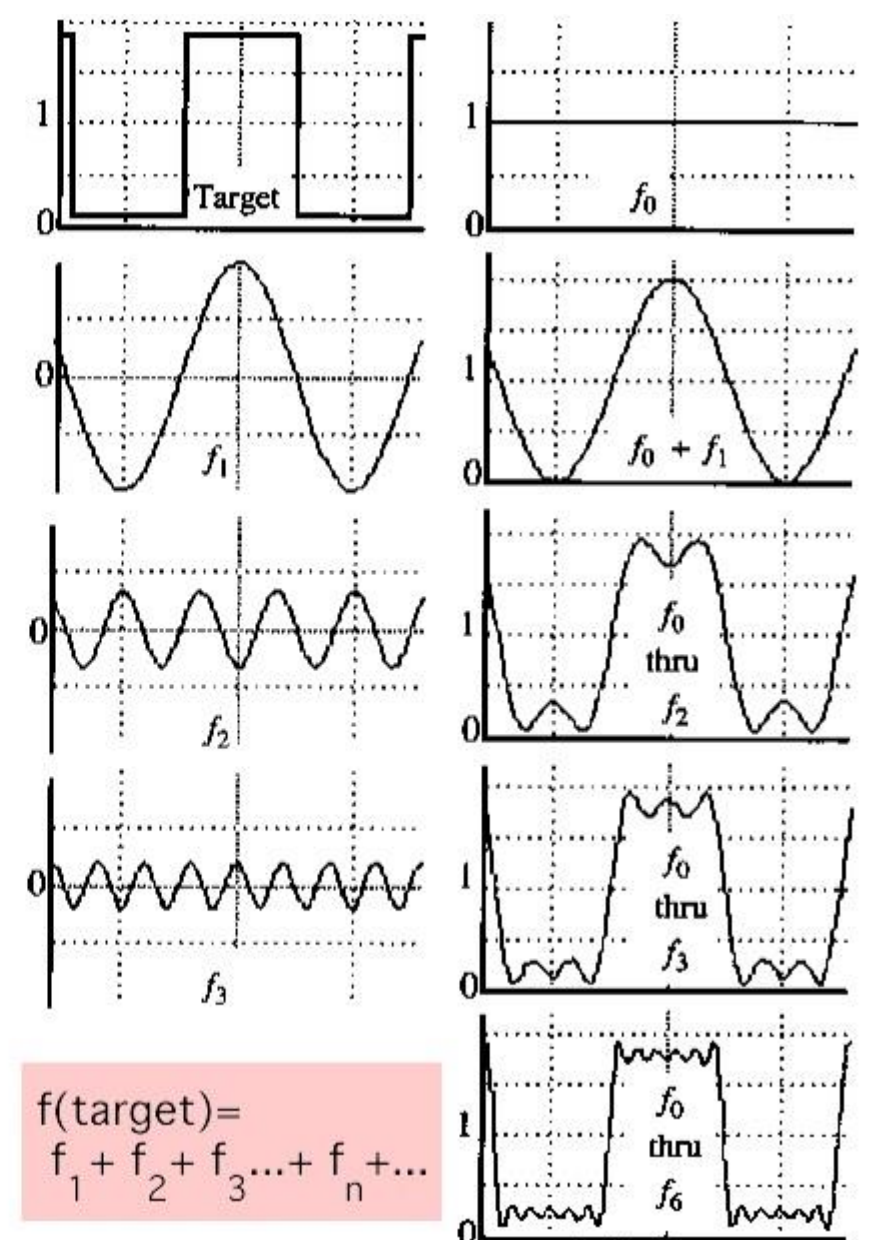


Image Fourier Analysis

- example : $g(t) = \sin(2\pi f t) + (1/3)\sin(2\pi(3f) t)$

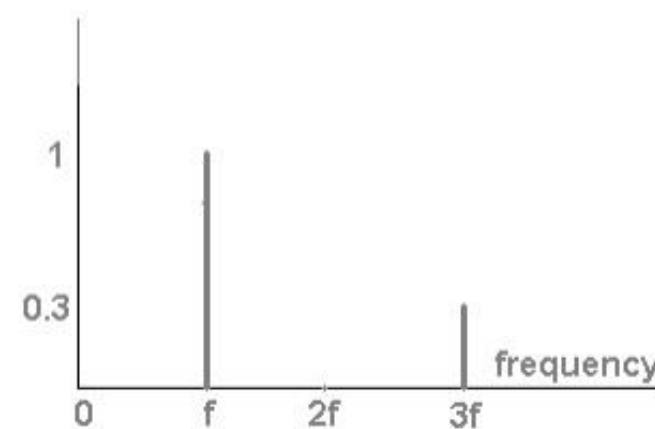
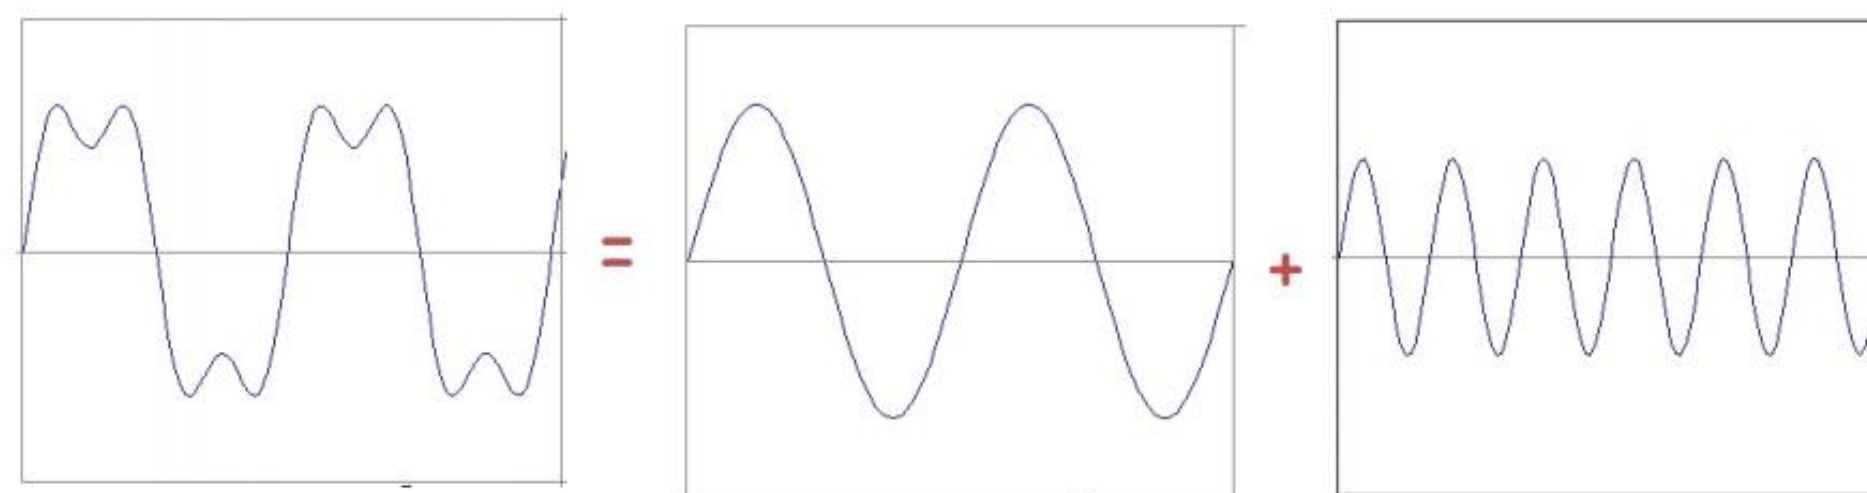


Image Fourier Analysis

Example: Music

- We think of music in terms of frequencies at different magnitudes

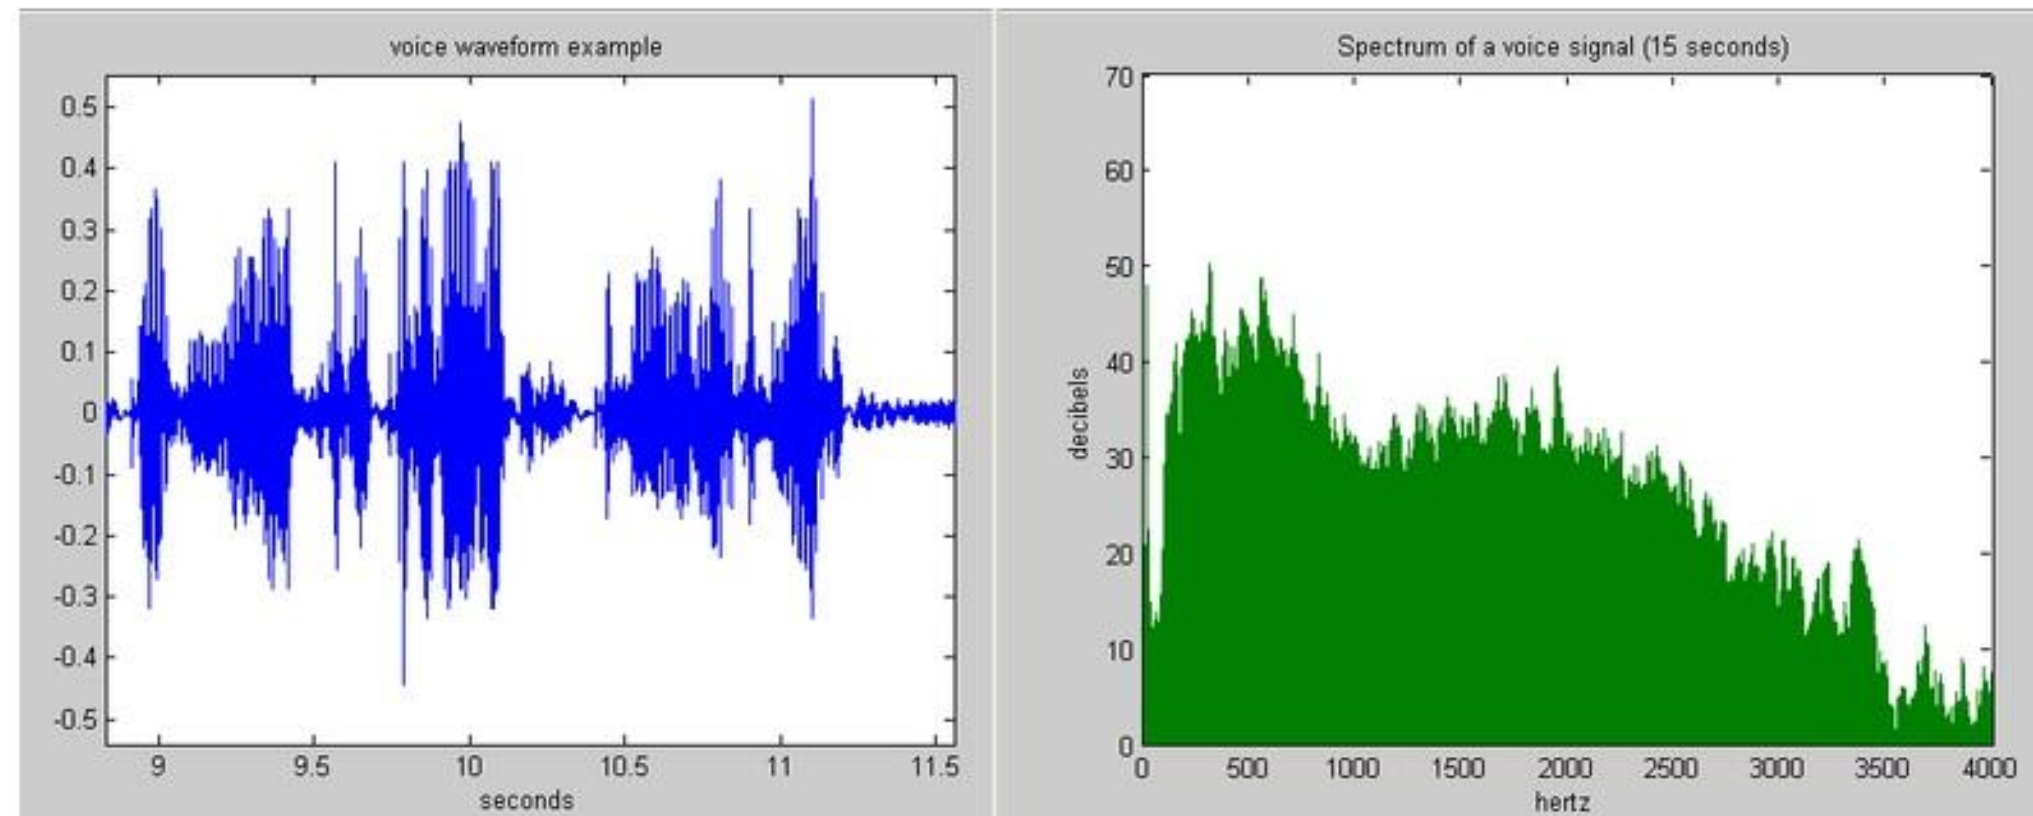
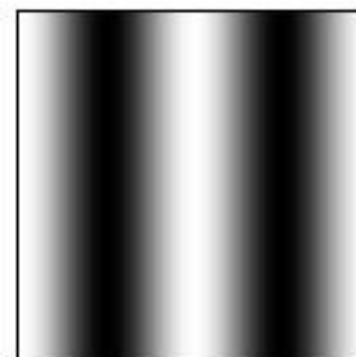


Image Fourier Analysis

Intensity Image



Fourier Image

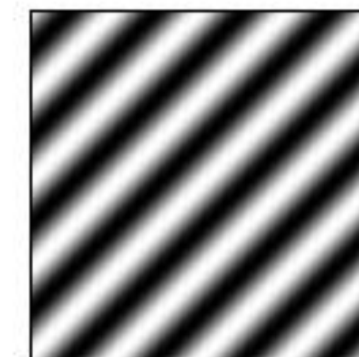
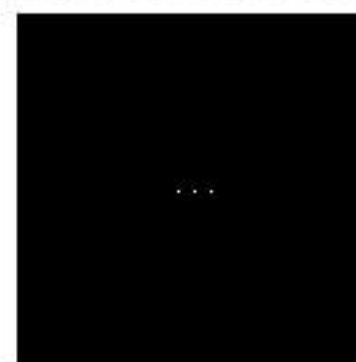
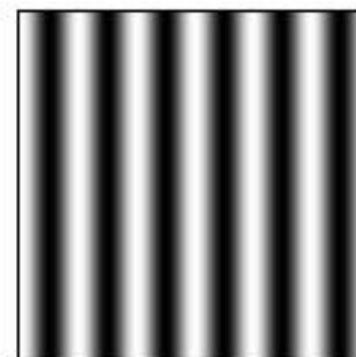
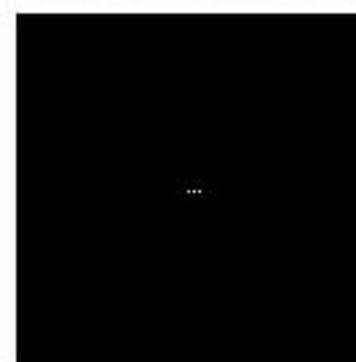


Image Fourier Analysis - Convolution

- The Fourier transform of the convolution of two functions is the product of their Fourier transforms

$$F[g * h] = F[g]F[h]$$

- **Convolution** in spatial domain is equivalent to **multiplication** in frequency domain!

$$g * h = F^{-1}[F[g]F[h]]$$

Image Fourier Analysis - Advantage

Too many spatial filter operations! What is the computational complexity of convolution?

Filtering in spatial domain

1	0	-1
2	0	-2
1	0	-1

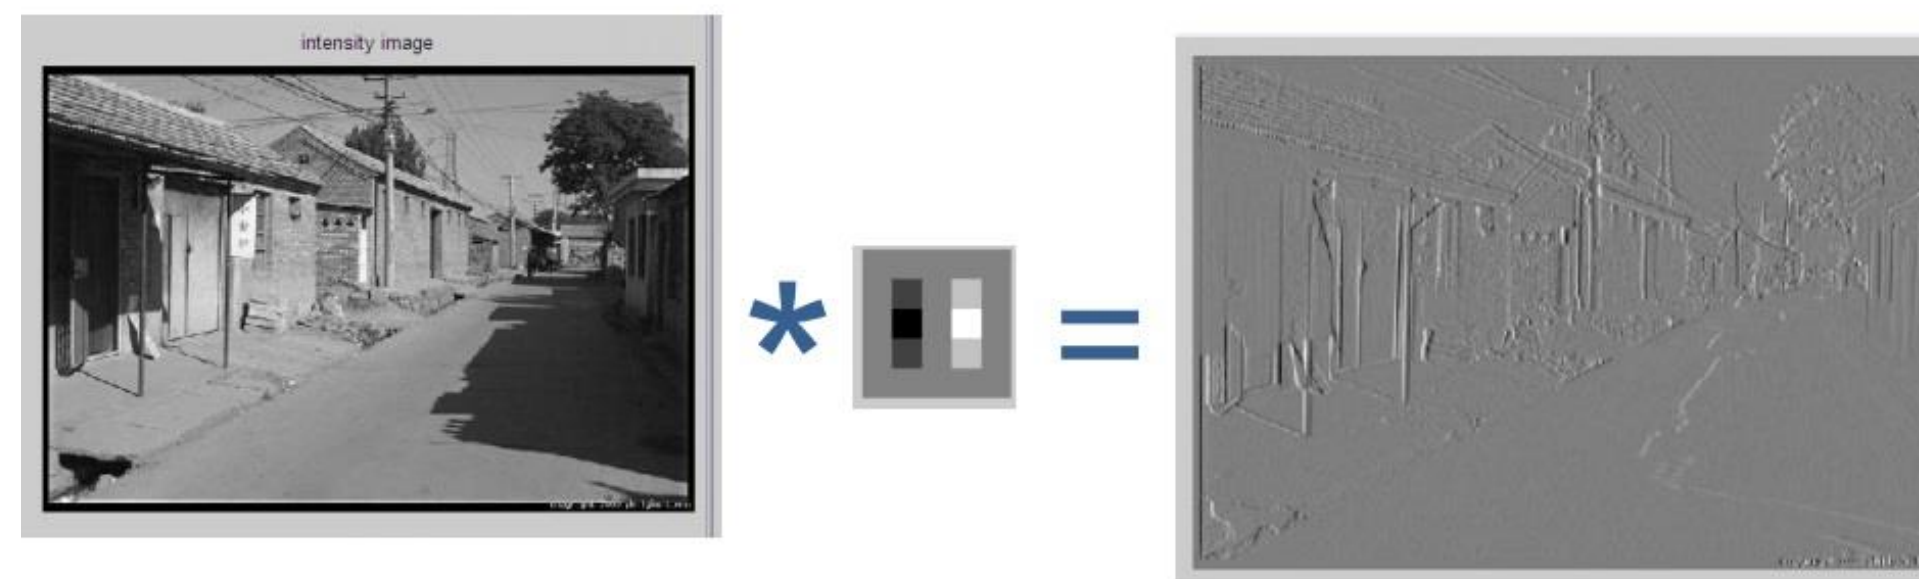


Image Fourier Analysis - Advantage

In frequency domain – Convolution becomes multiplication

Filtering in frequency domain

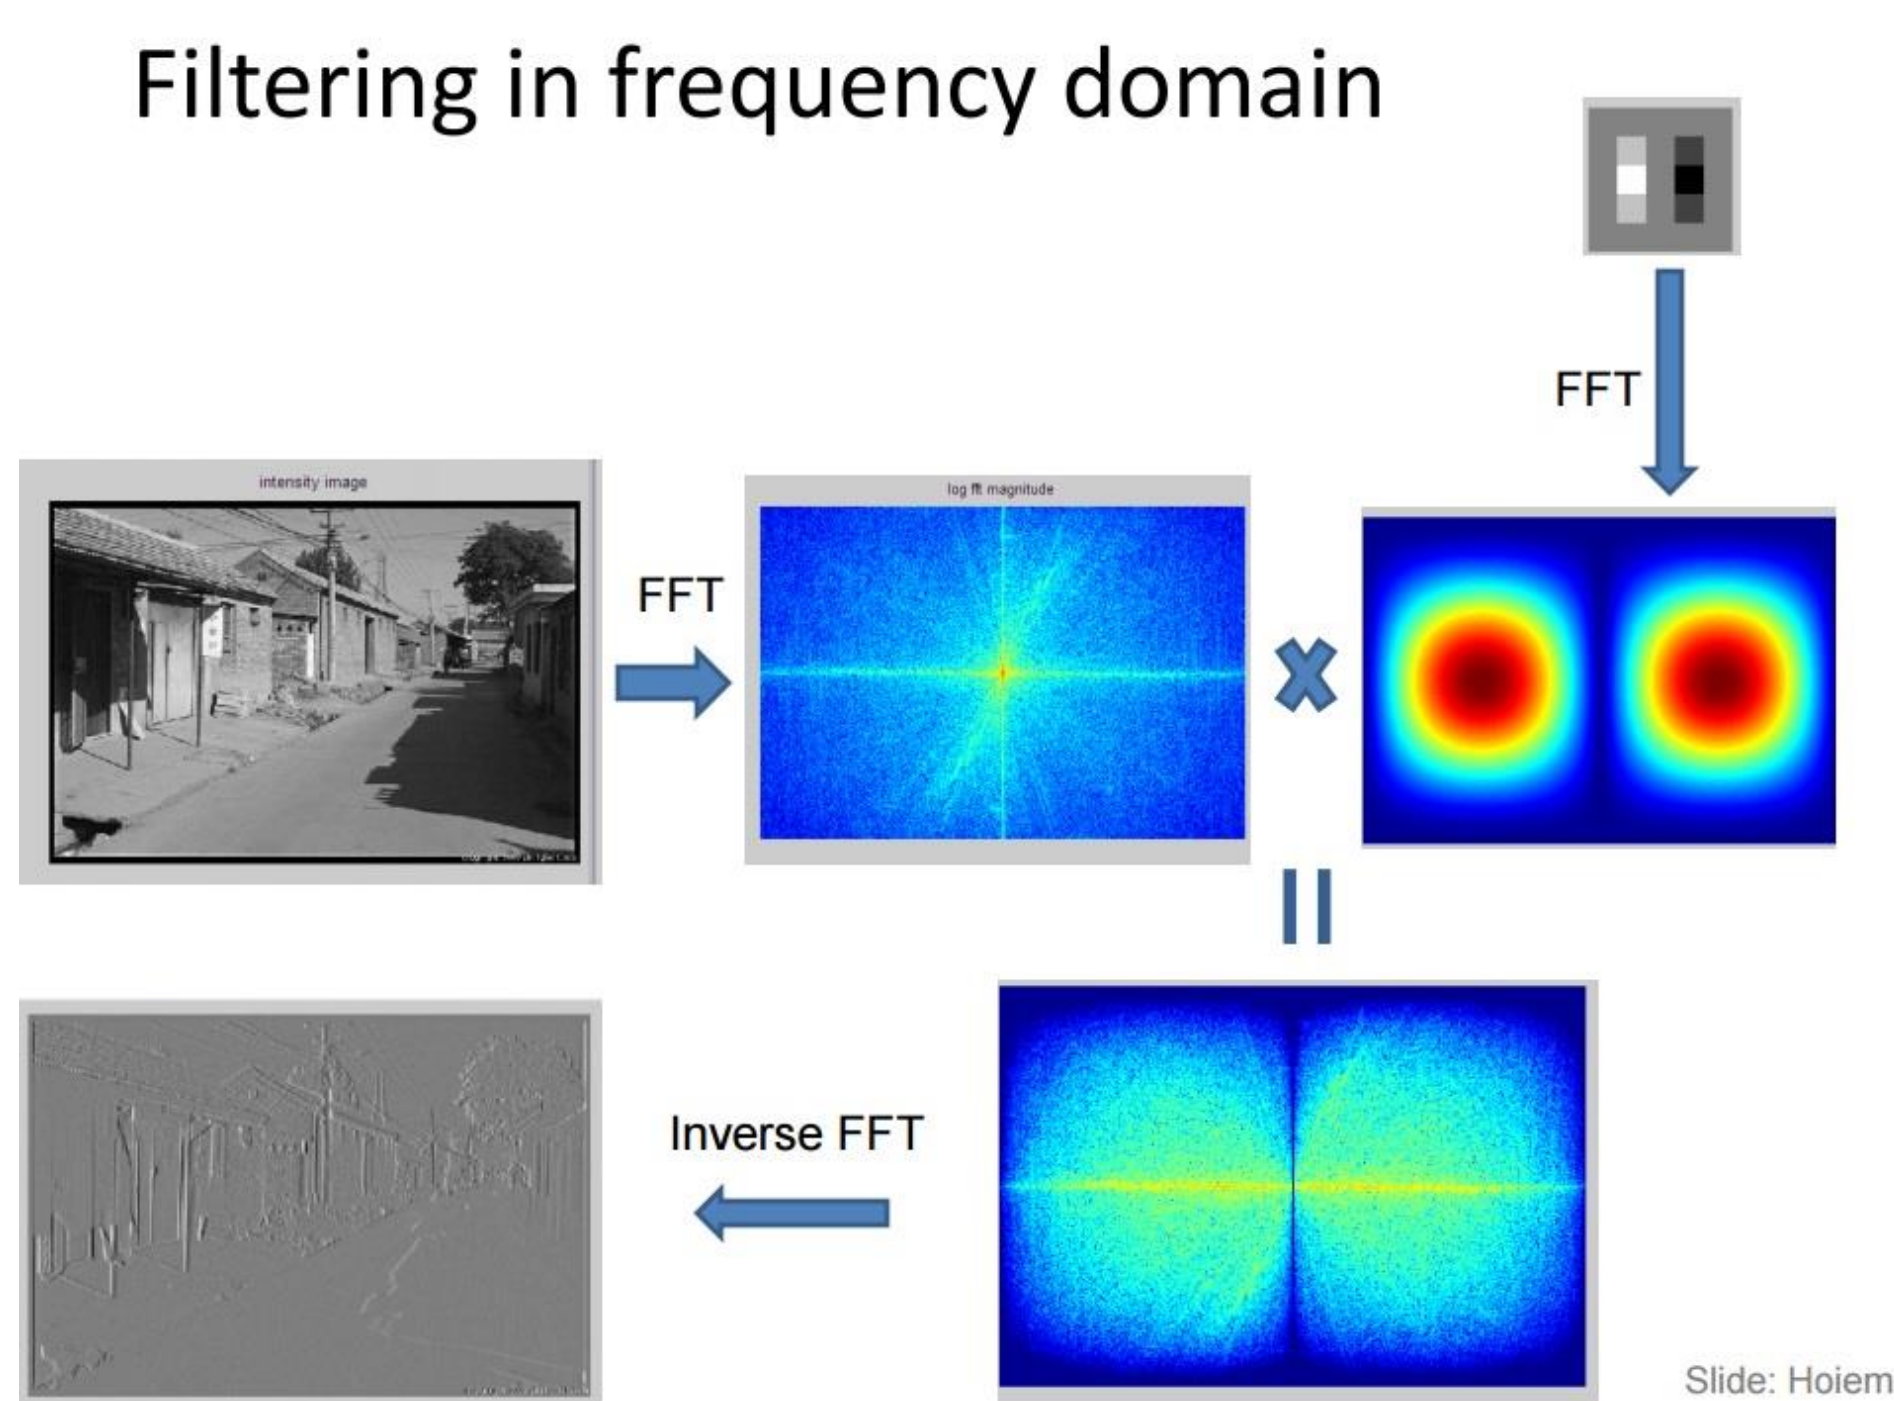
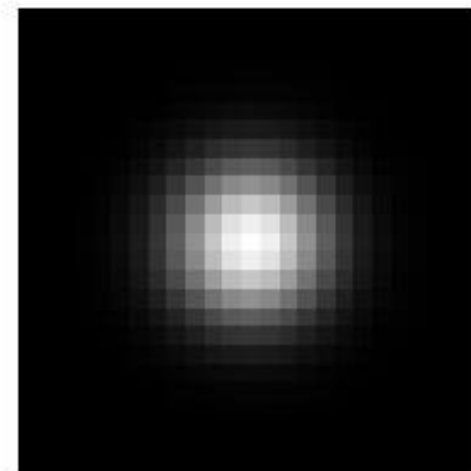
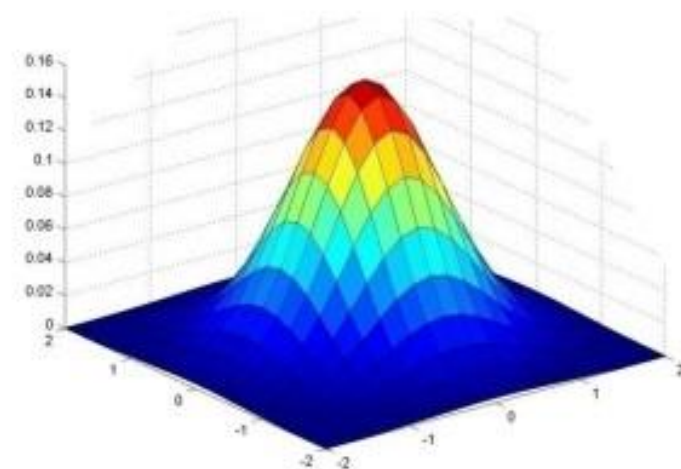


Image Smoothing – Gaussian Kernel



$$G_{\sigma} = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$

Image Smoothing – Gaussian Kernel

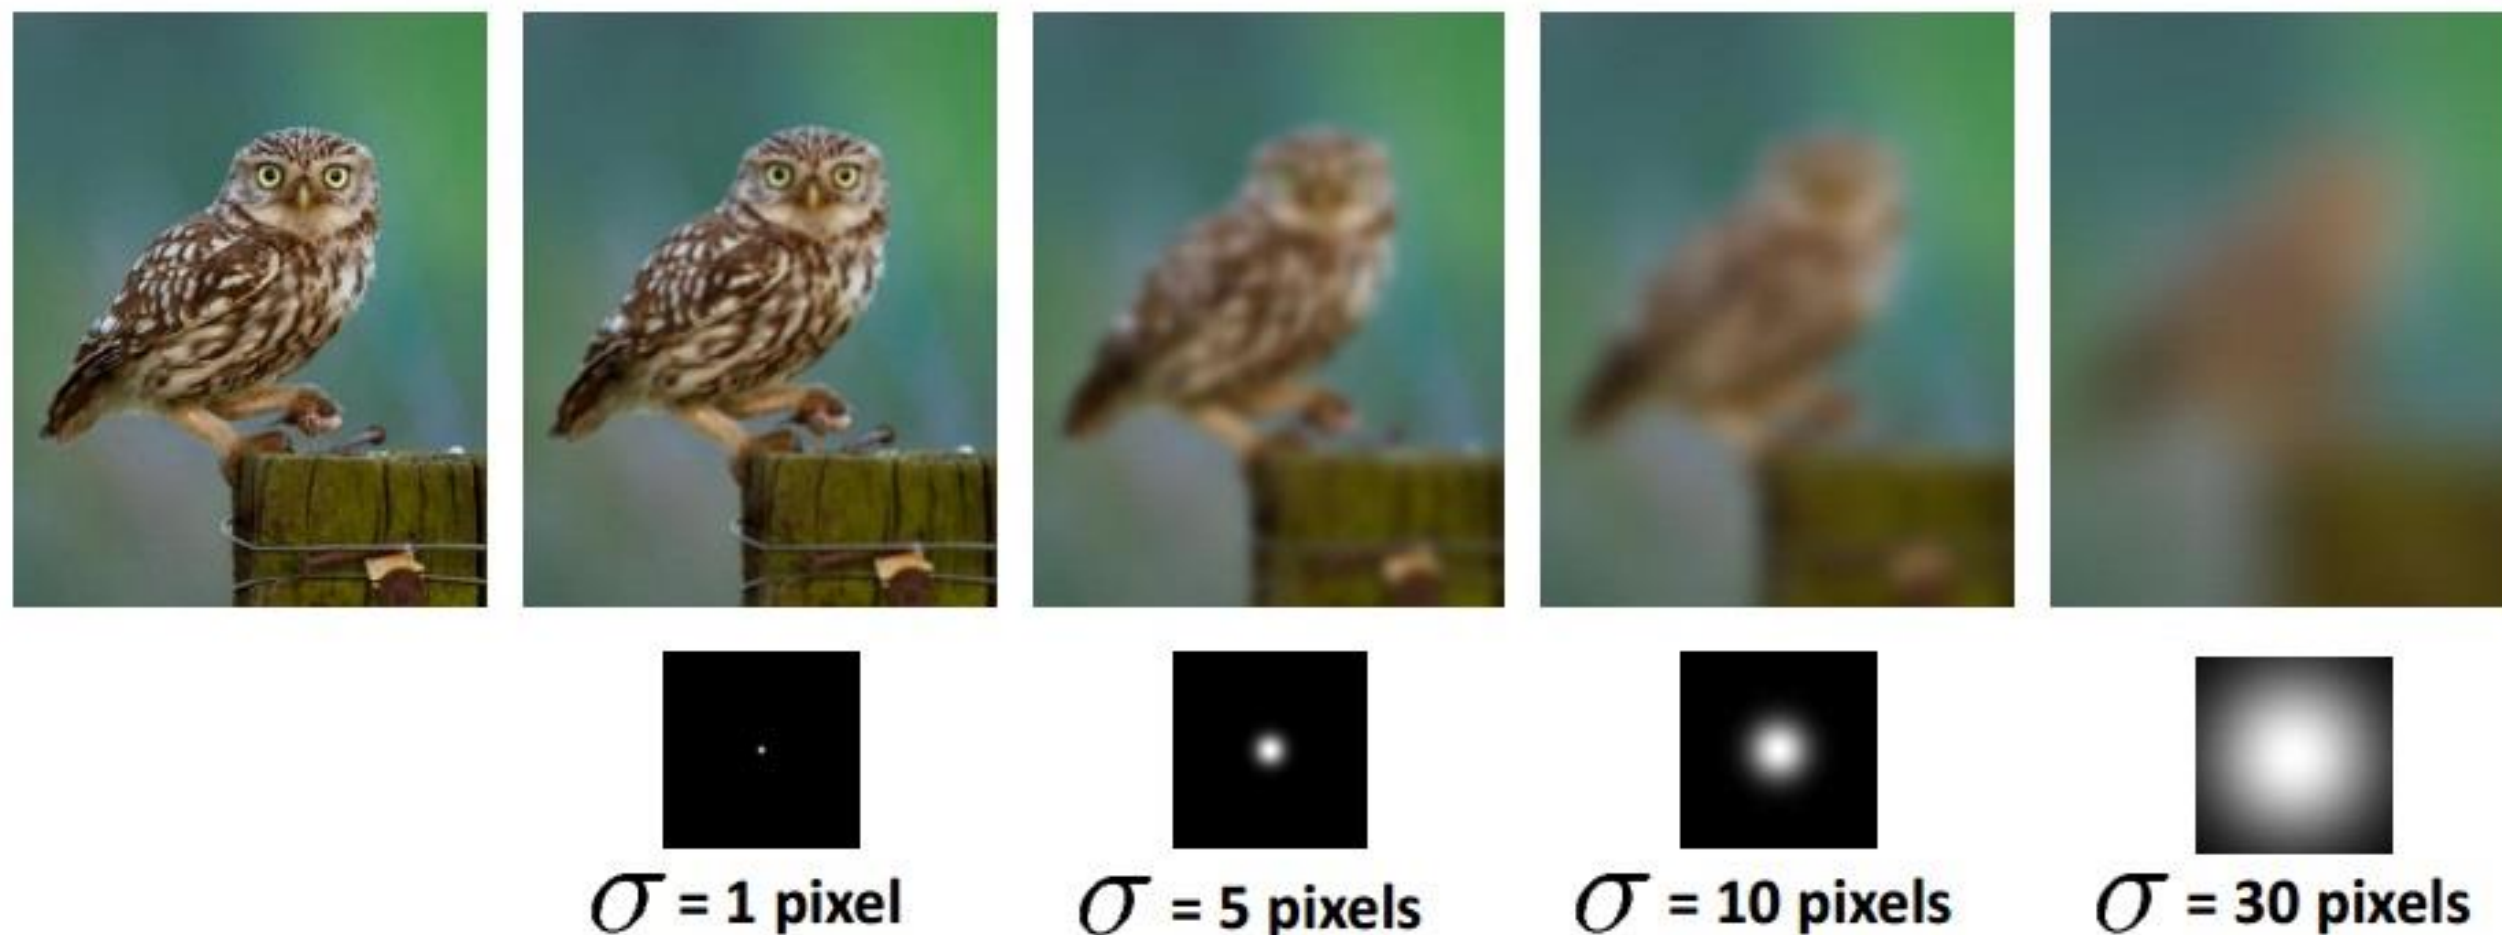


Image Smoothing – Gaussian Kernel

- Removes “high-frequency” components from the image (low-pass filter)
- Convolution with self is another Gaussian

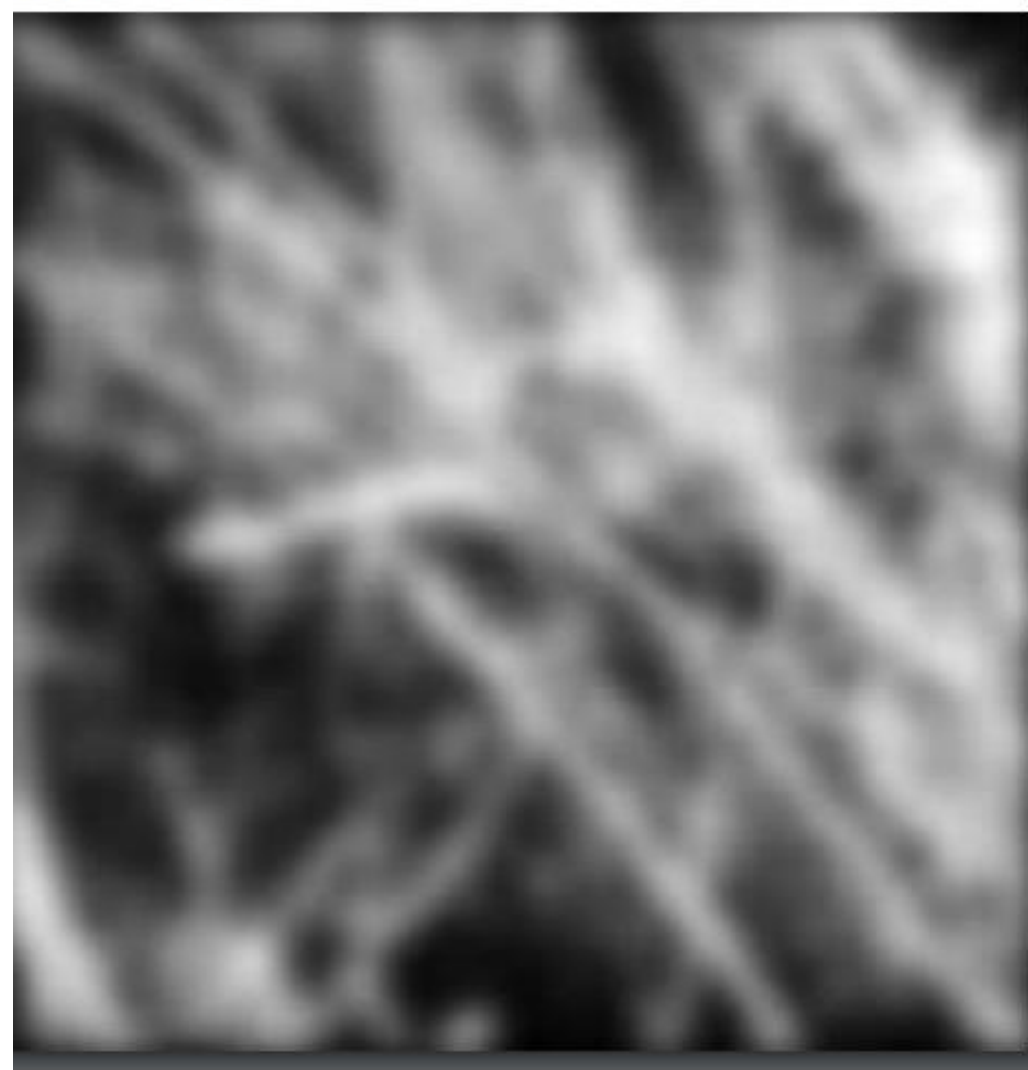


- Convolving two times with Gaussian kernel of width σ = convolving once with kernel of width $\sigma\sqrt{2}$

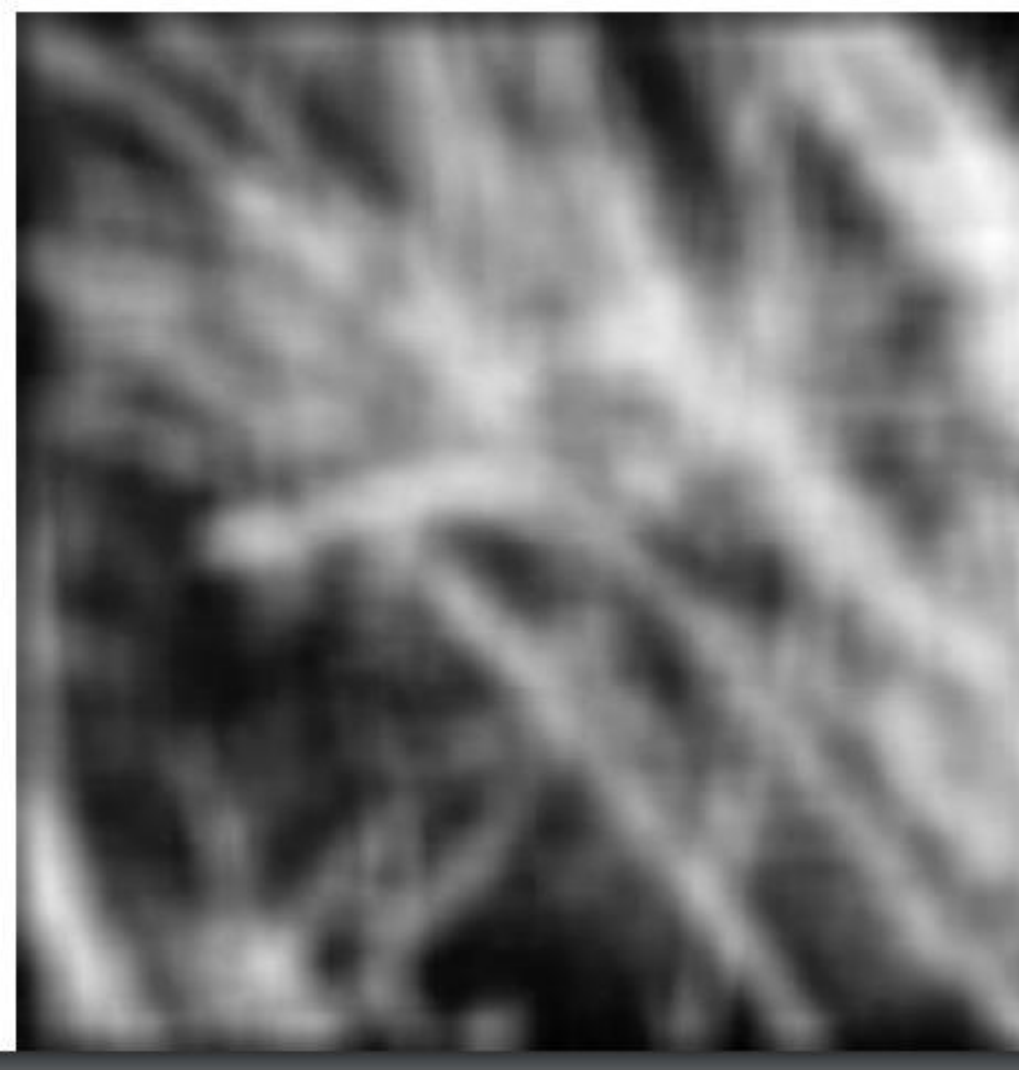
Box Filter Vs Gaussian Filter

Why does the Gaussian give a nice smooth image, but the square filter give edgy artifacts?

Gaussian

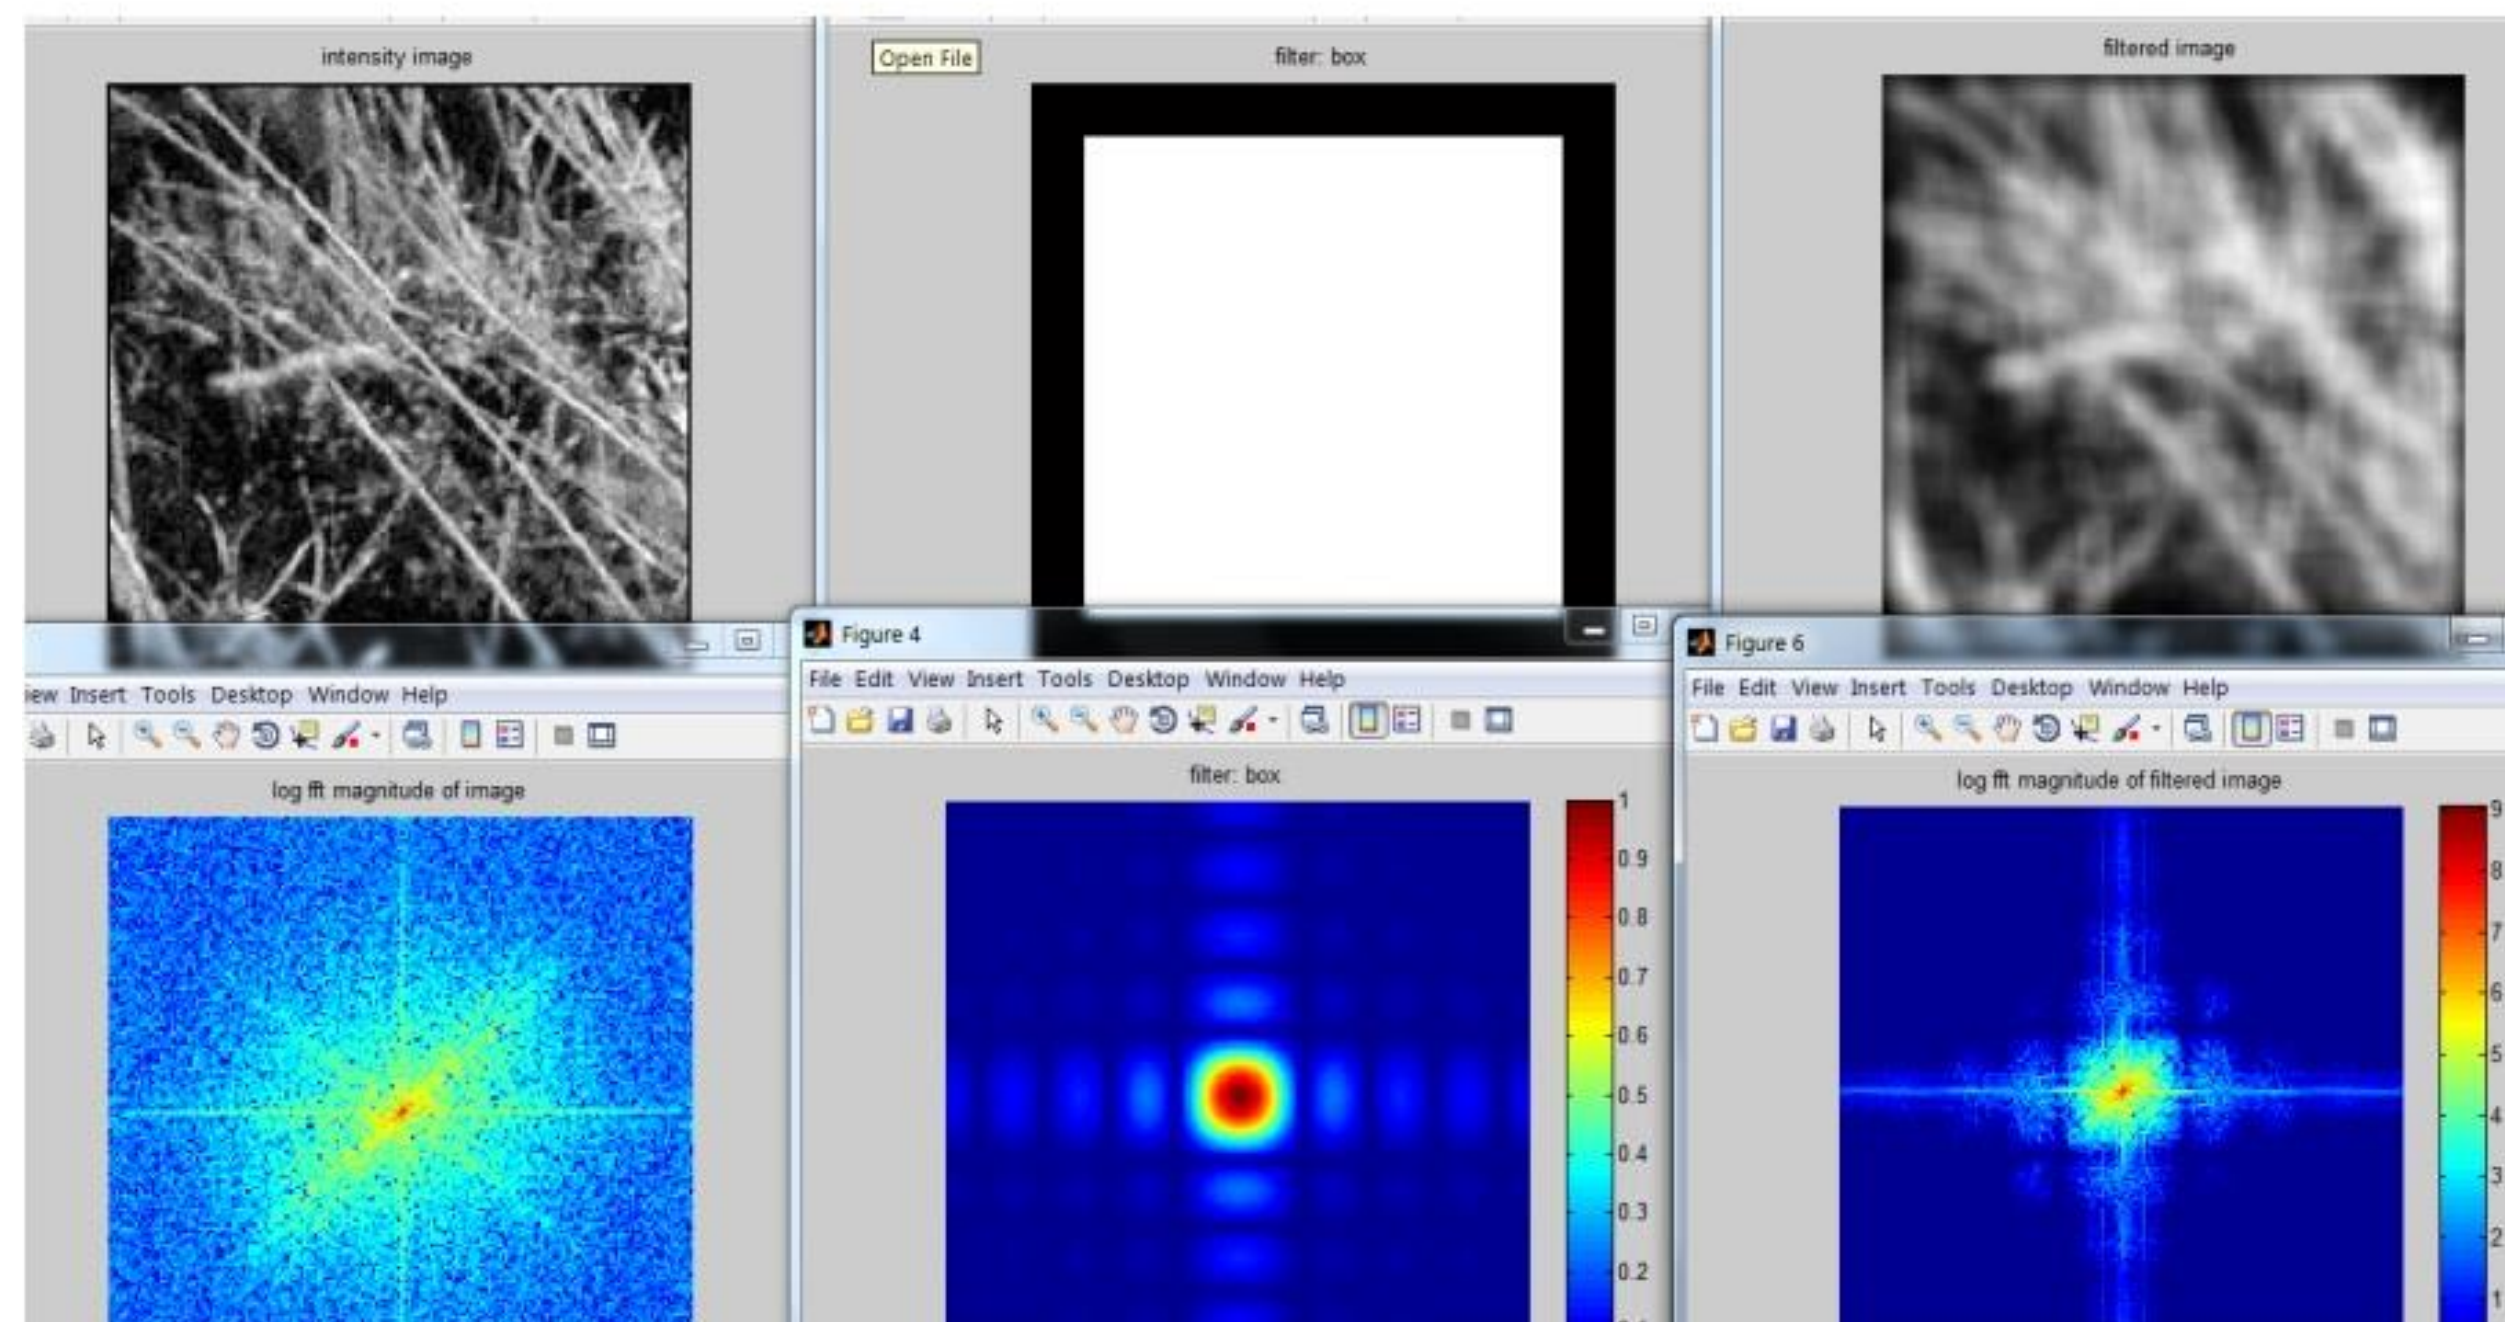


Box filter



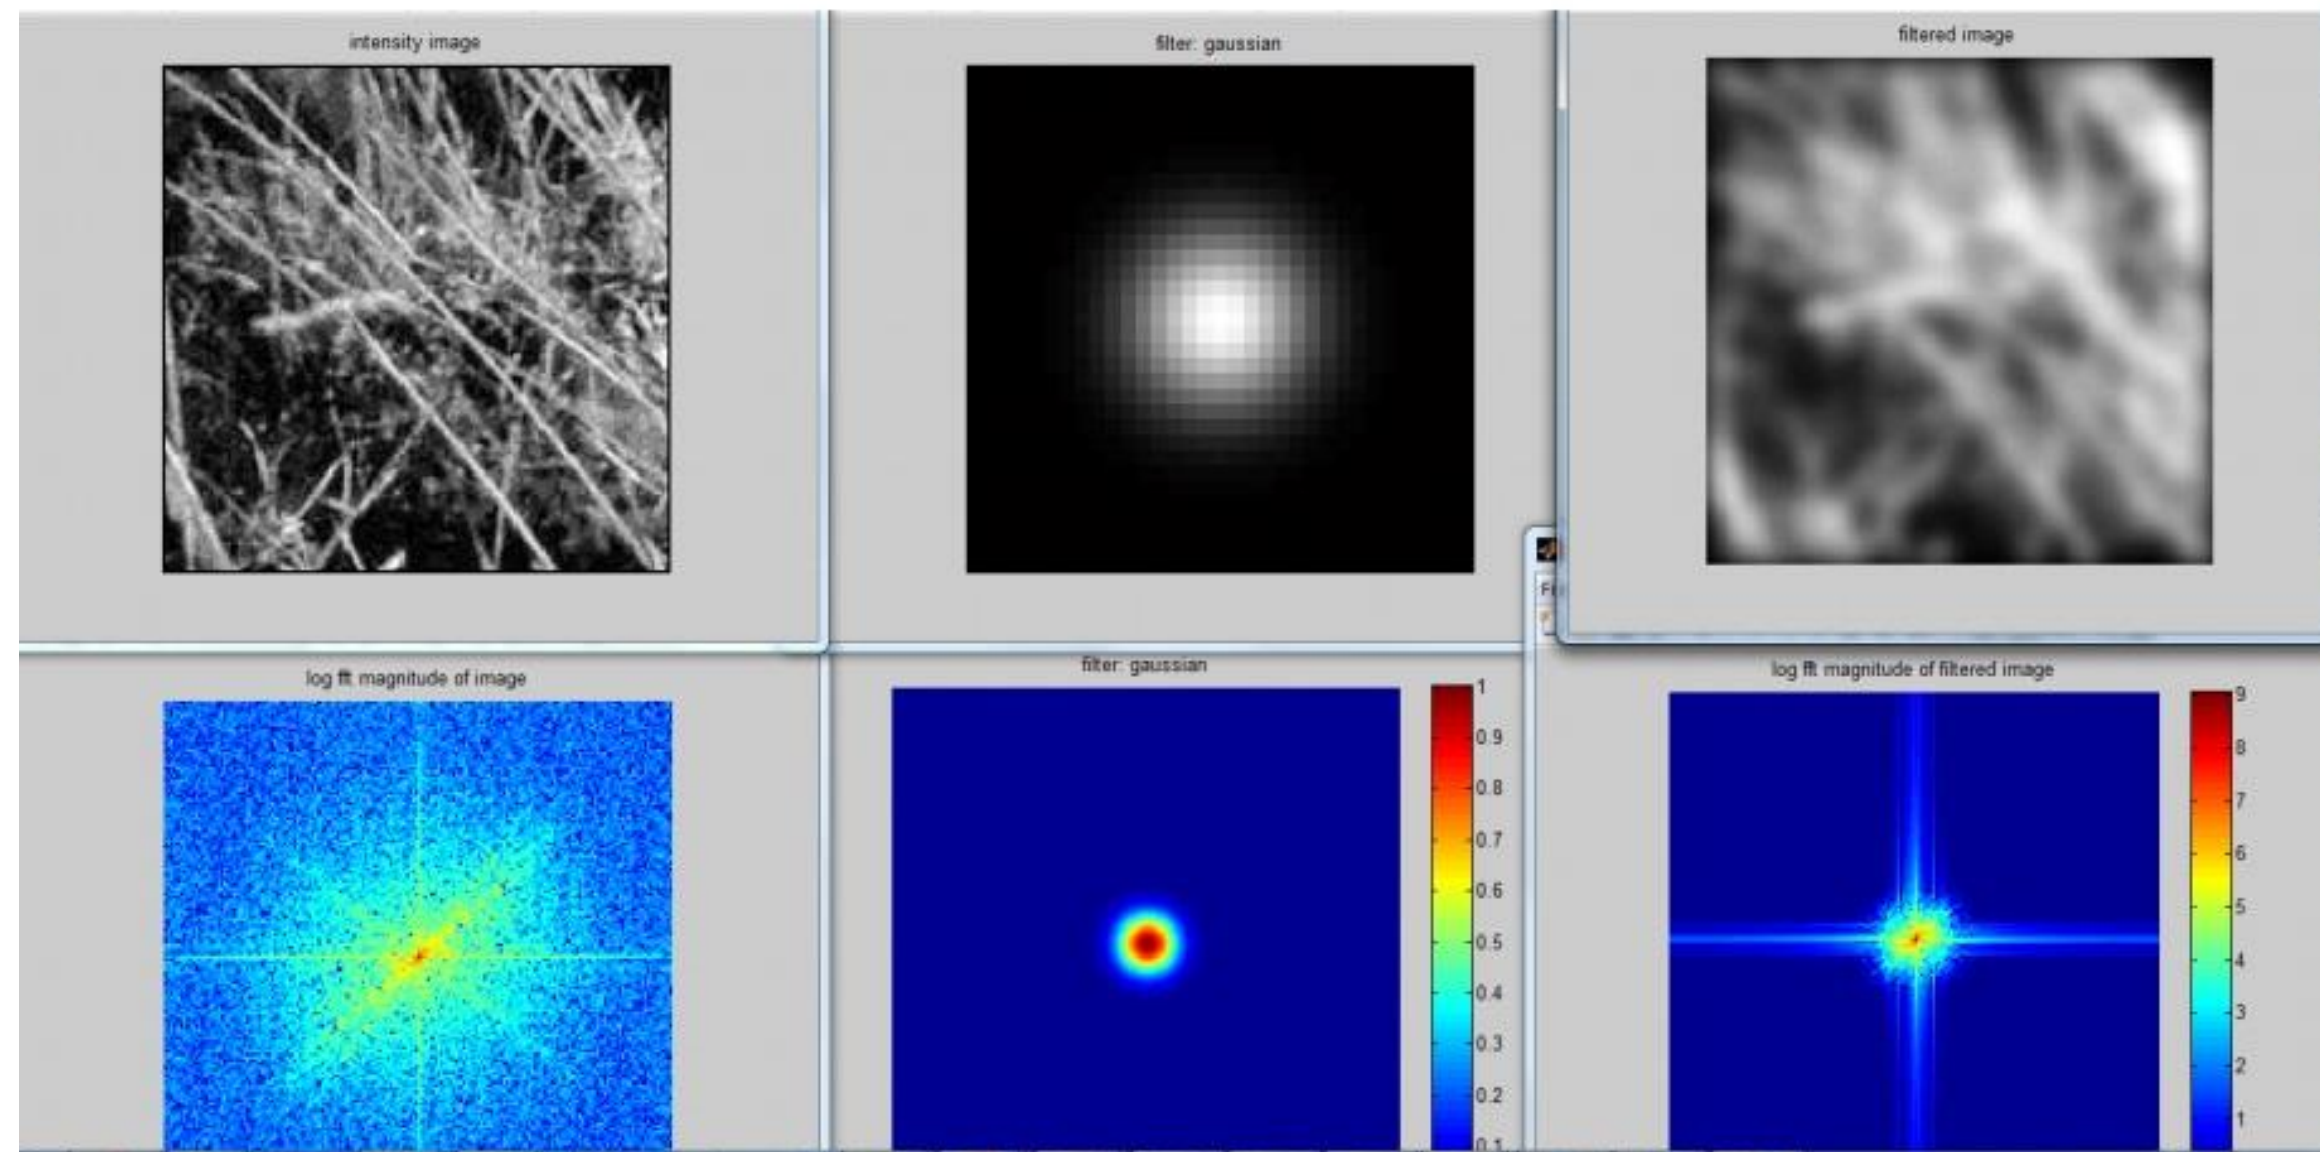
Fourier Analysis – Box Filter

Box Filter




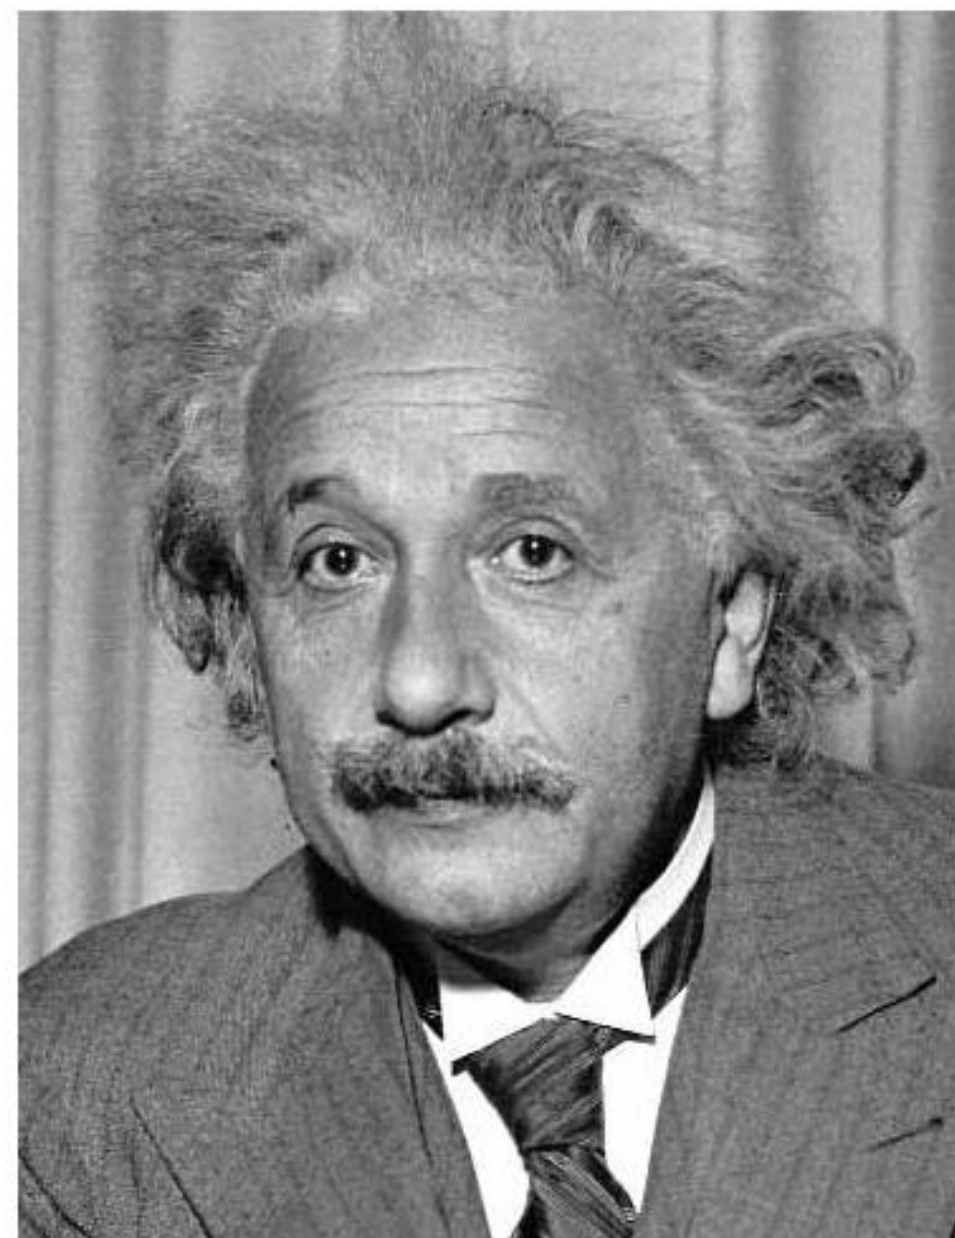
Fourier Analysis – Gaussian Filter

Gaussian



Template Matching

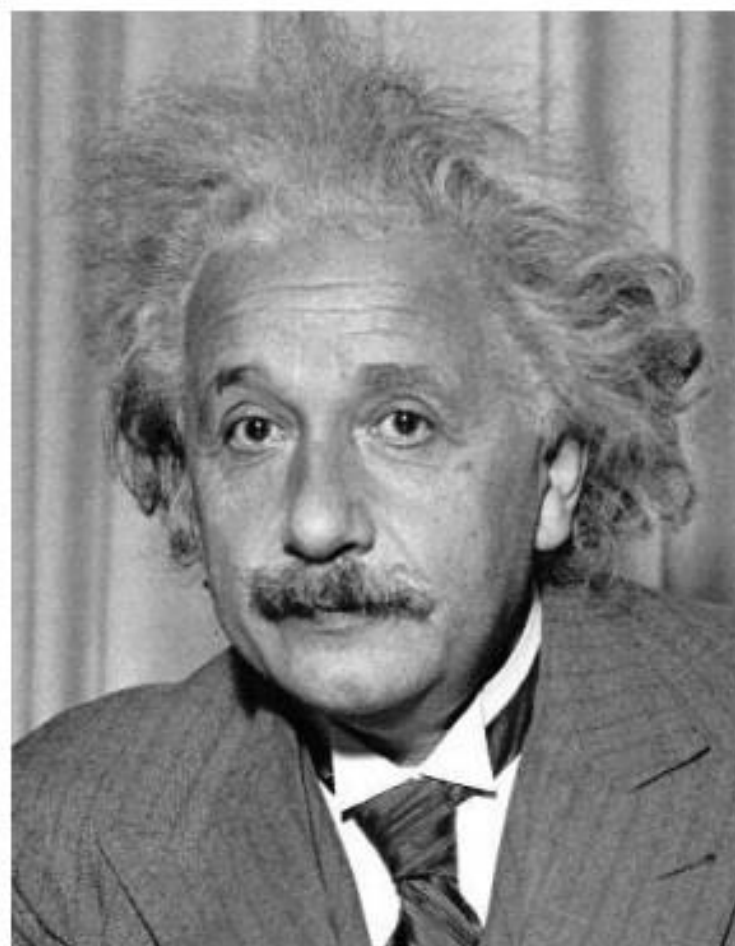
- Goal: find  in image
- Main challenge: What is a good similarity or distance measure between two patches?
 - Correlation
 - Zero-mean correlation
 - Sum Square Difference
 - Normalized Cross Correlation



Sum of Squared Differences (SSD)

- Goal: find  in image
- Method 2: SSD

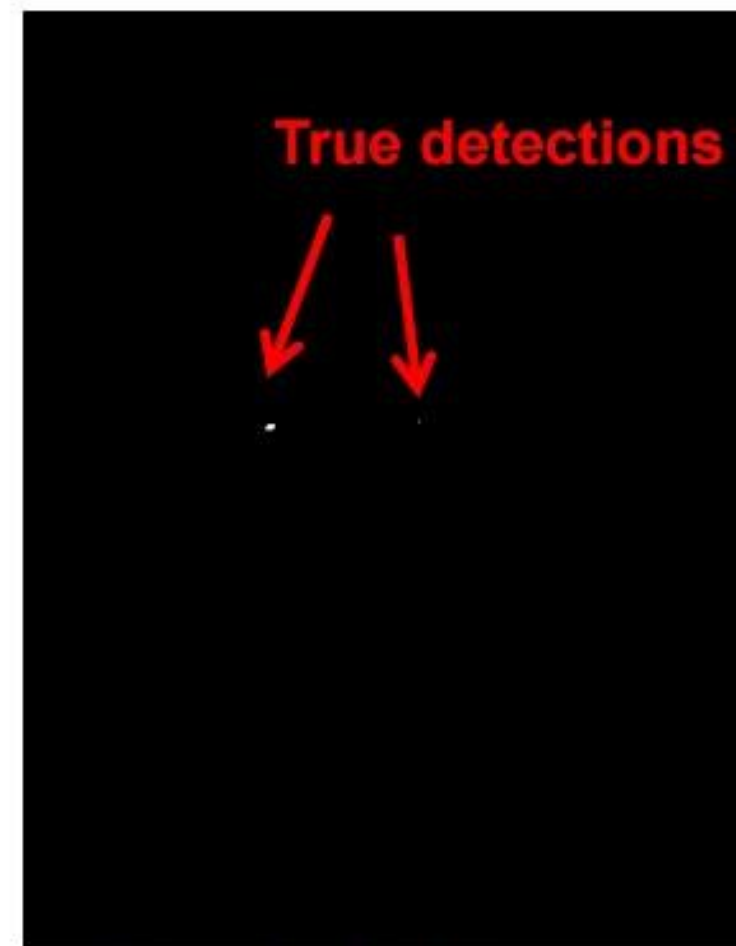
$$h[m,n] = \sum_{k,l} (g[k,l] - f[m+k, n+l])^2$$



Input



1- sqrt(SSD)



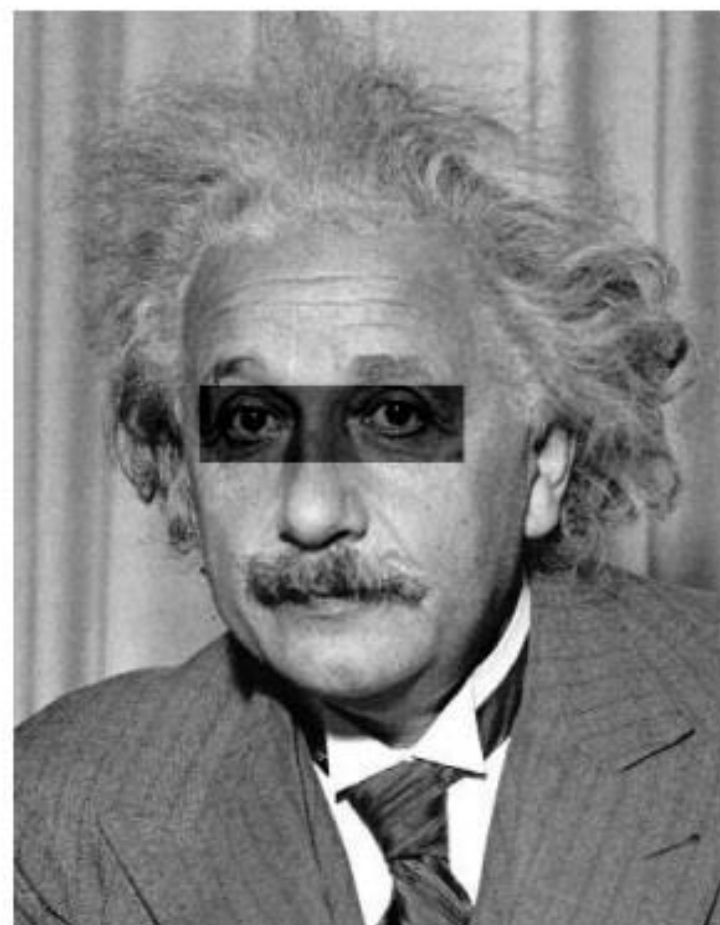
Thresholded Image

Sum of Squared Differences (SSD)

- Goal: find  in image
- Method 2: SSD

What's the potential
downside of SSD?

$$h[m,n] = \sum_{k,l} (g[k,l] - f[m+k,n+l])^2$$



Input



1- sqrt(SSD)

Normalized Cross Correlation

- Goal: find  in image
- Method 3: Normalized cross-correlation

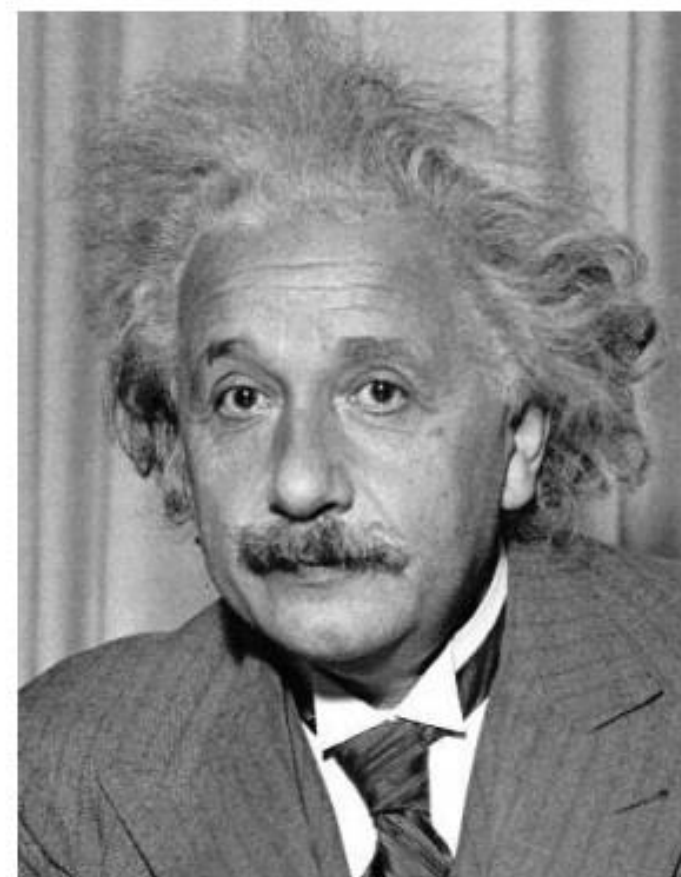
$$h[m,n] = \frac{\sum_{k,l} (g[k,l] - \bar{g})(f[m-k, n-l] - \bar{f}_{m,n})}{\left(\sum_{k,l} (g[k,l] - \bar{g})^2 \sum_{k,l} (f[m-k, n-l] - \bar{f}_{m,n})^2 \right)^{0.5}}$$

mean template mean image patch

↓ ↓

Normalized Cross Correlation

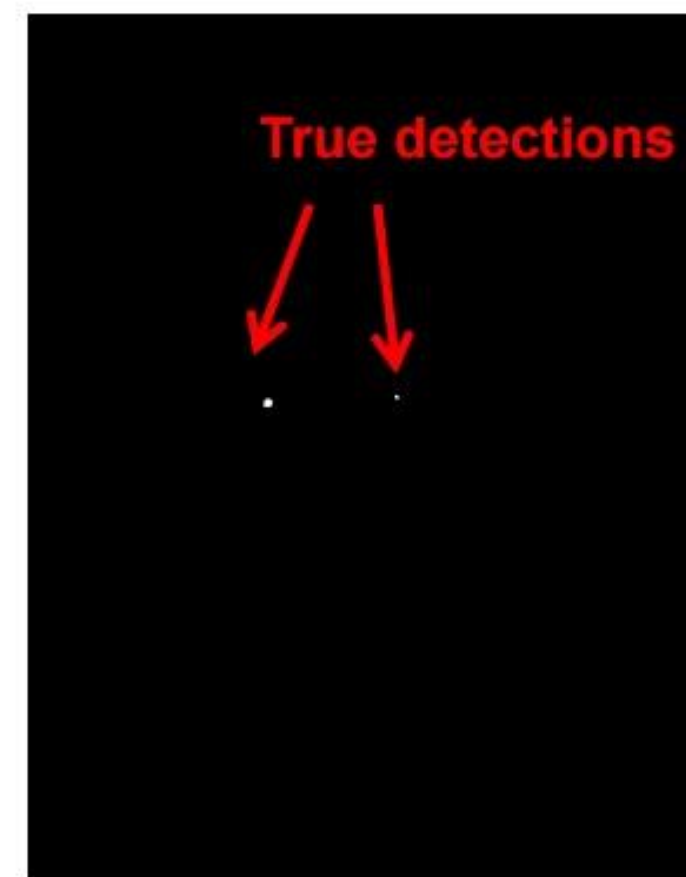
- Goal: find  in image
- Method 3: Normalized cross-correlation



Input



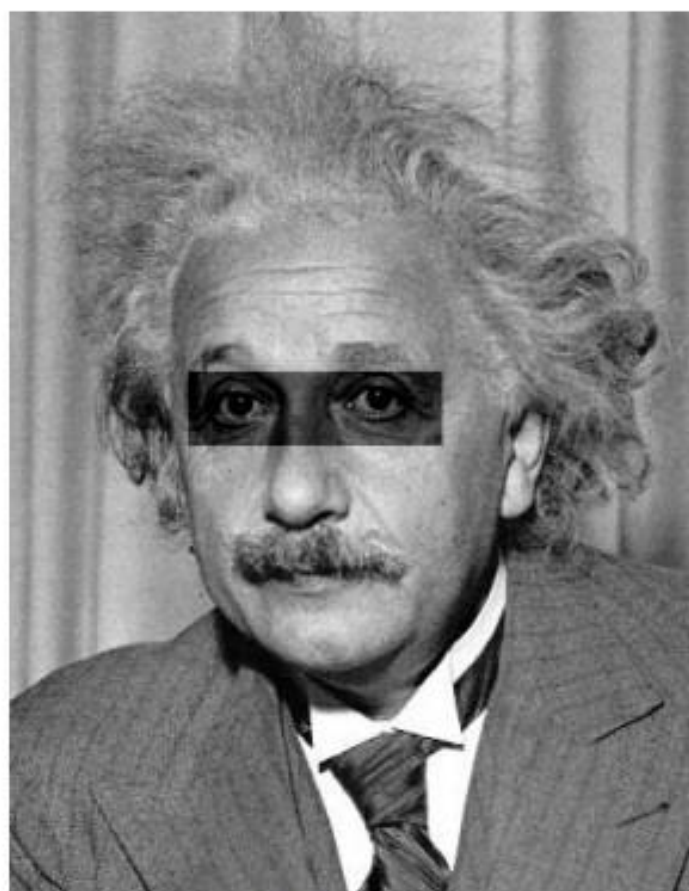
Normalized X-Correlation



Thresholded Image

Normalized Cross Correlation

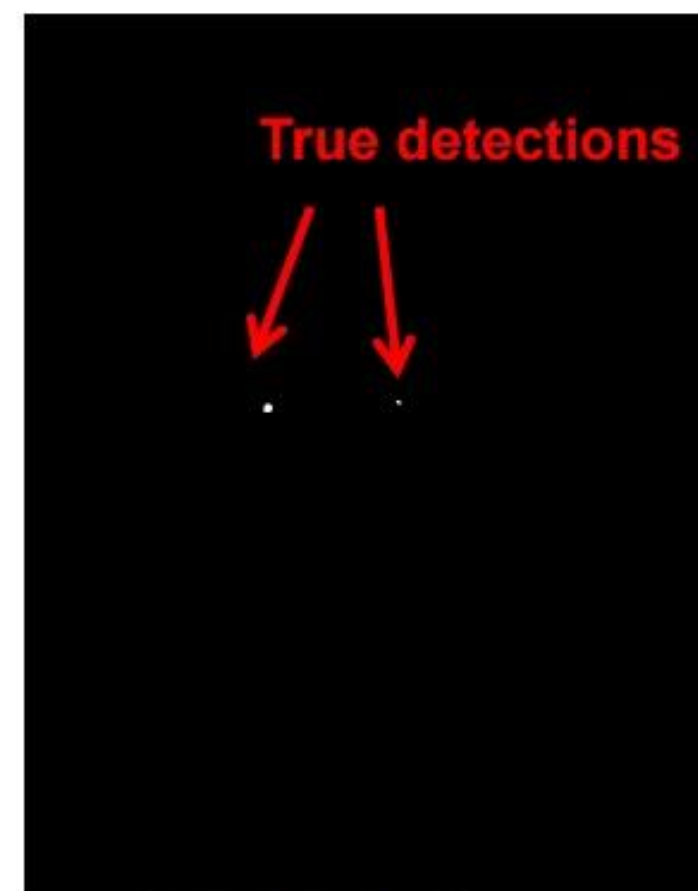
- Goal: find  in image
- Method 3: Normalized cross-correlation



Input



Normalized X-Correlation



Thresholded Image

SSD vs NCC – What to use?

A: Depends

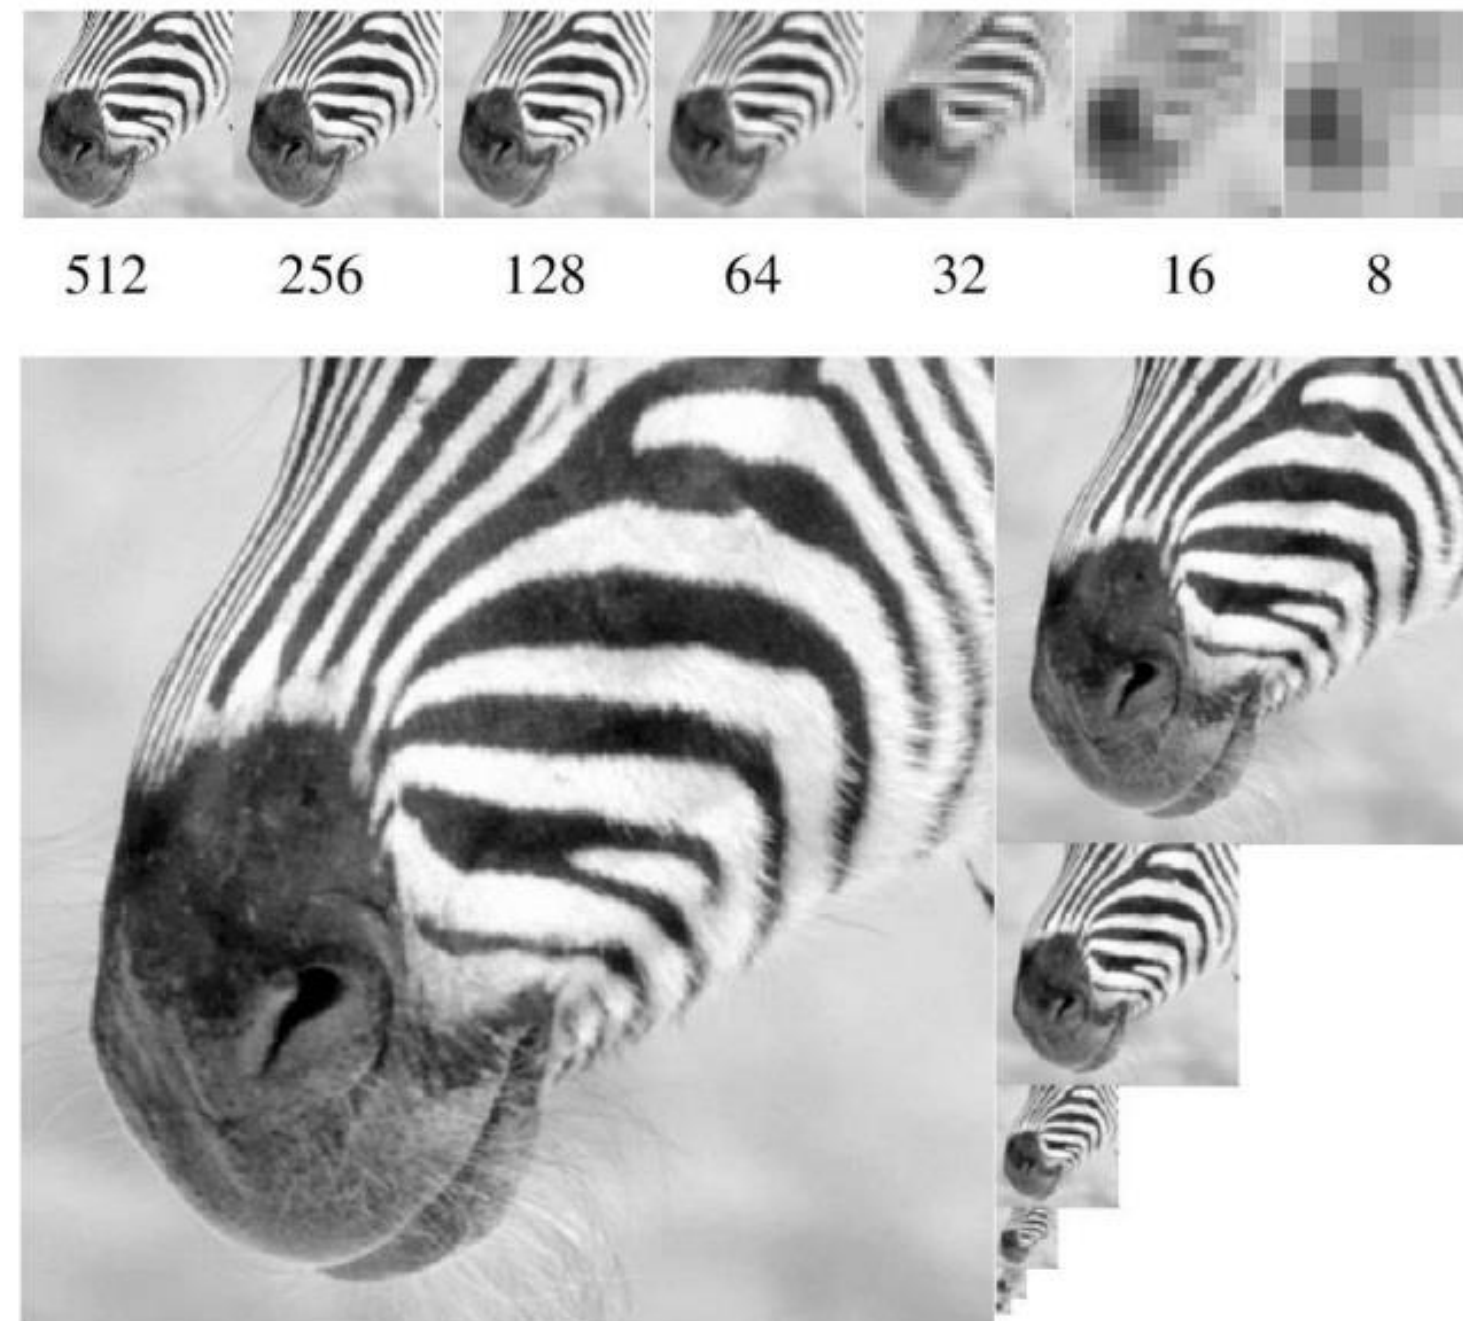
- SSD: faster, sensitive to overall intensity
- Normalized cross-correlation: slower, invariant to local average intensity and contrast

Image Pyramids

Q: What if we want to find larger or smaller eyes?

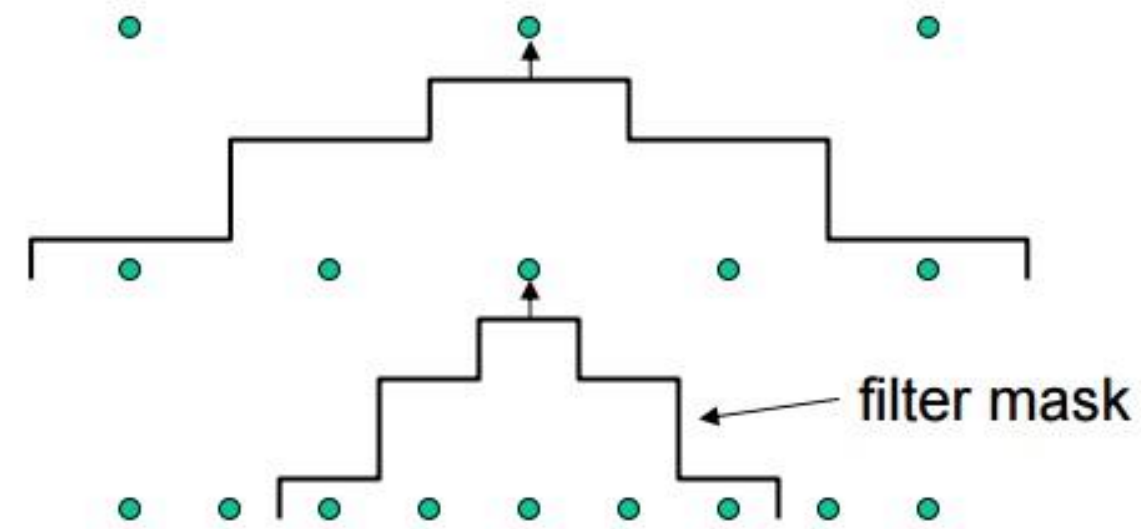
A: Image Pyramid

Image Pyramids



Source: Forsyth

Image Pyramids



Repeat

- Filter
- Subsample

Until minimum resolution reached

- can specify desired number of levels (e.g., 3-level pyramid)

The whole pyramid is only $\frac{4}{3}$ the size of the original image!

Gaussian Smoothing

What does blurring take away?



original

Gaussian Smoothing

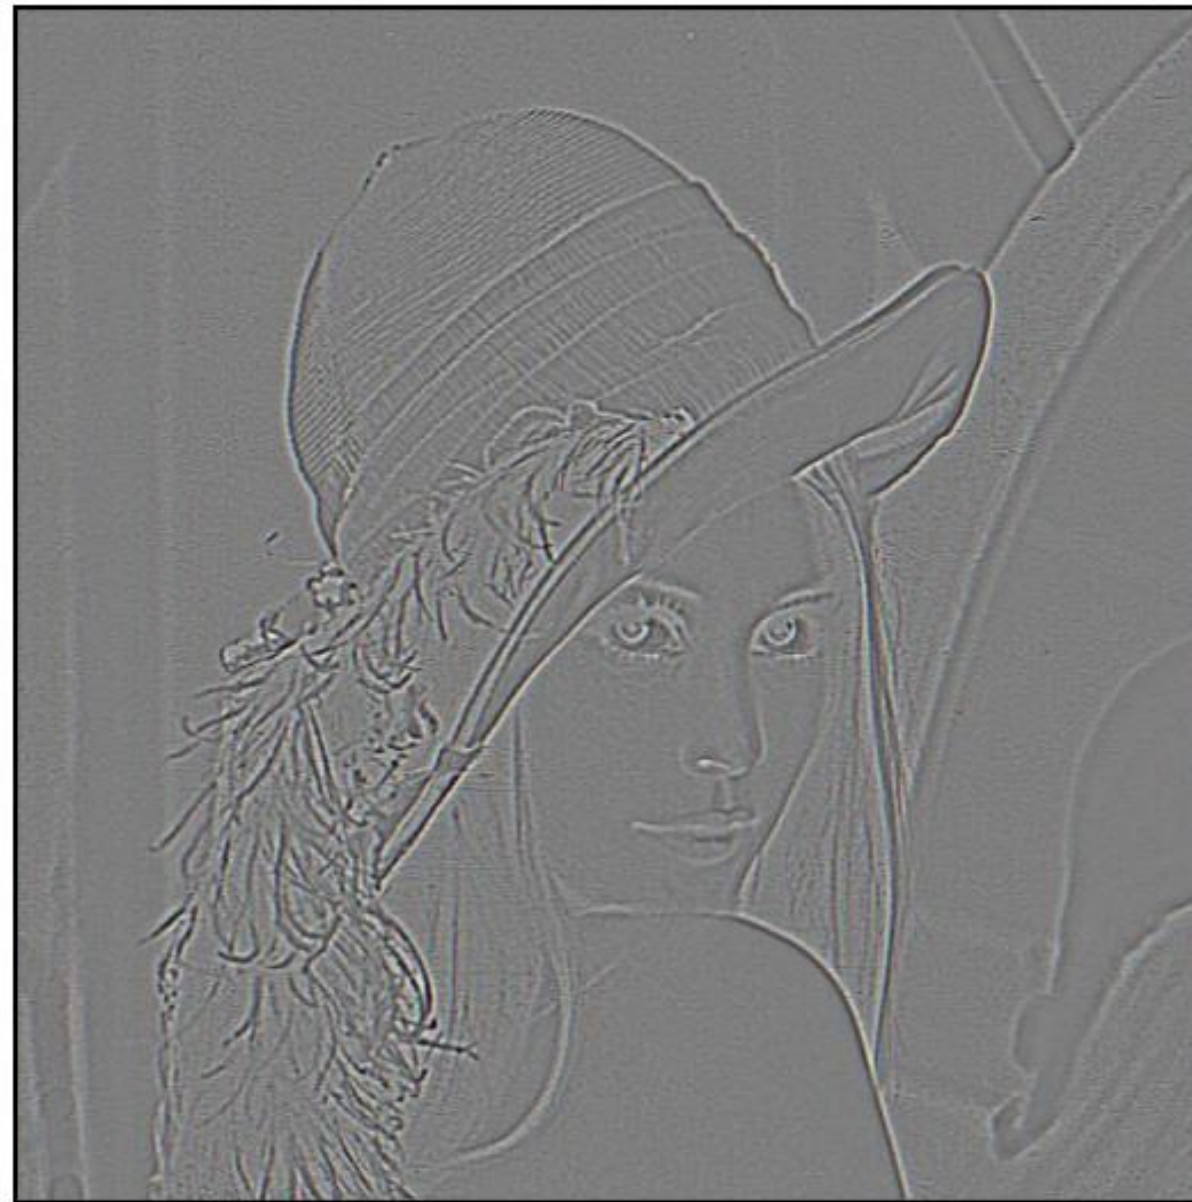
What does blurring take away?



smoothed (5x5 Gaussian)

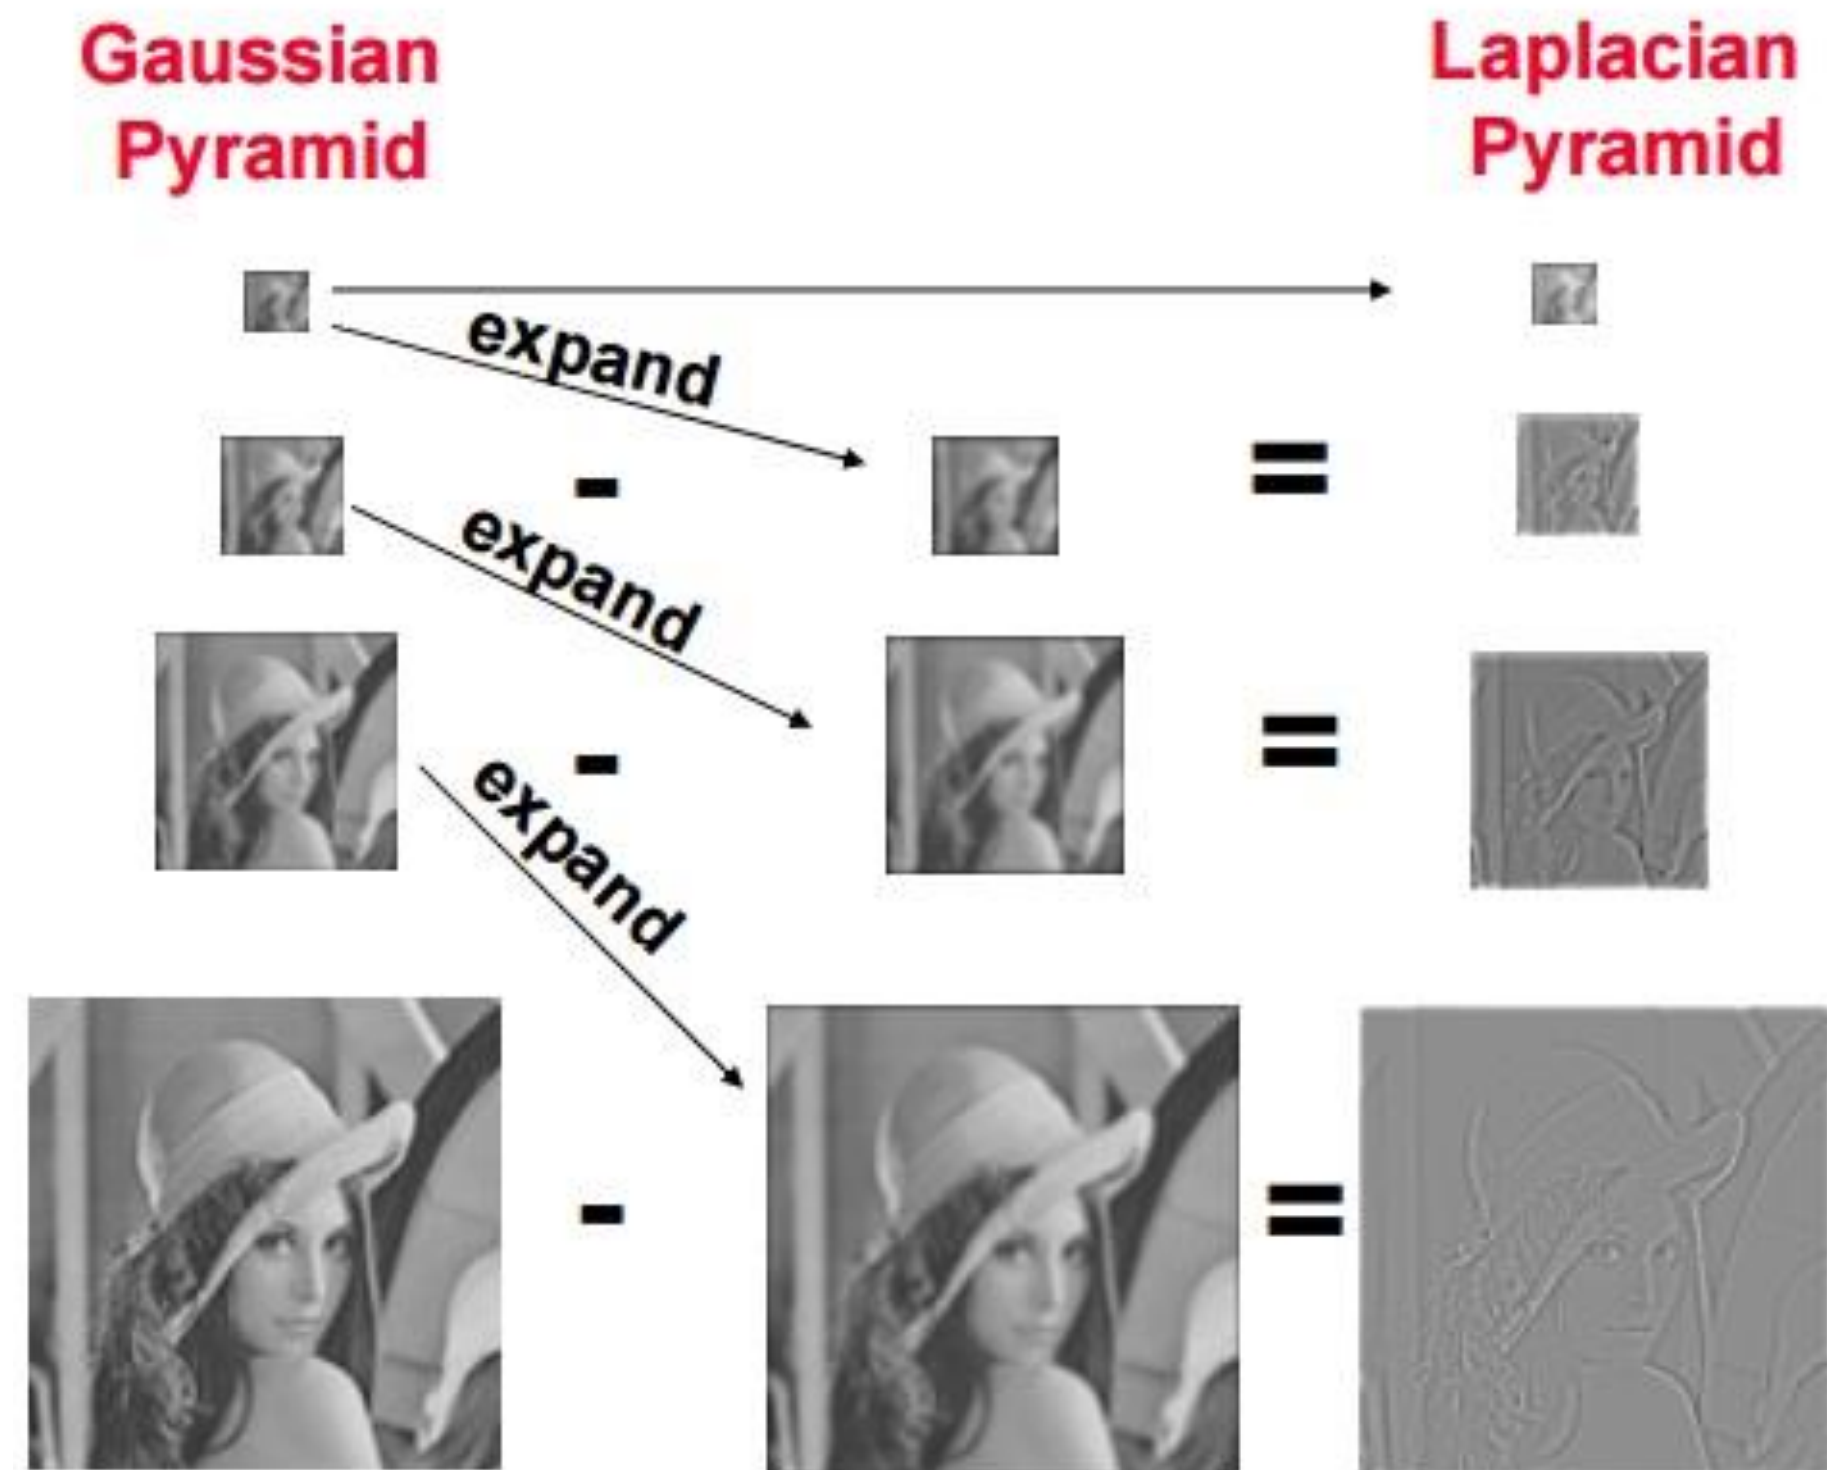
Gaussian Smoothing

High-Pass filter

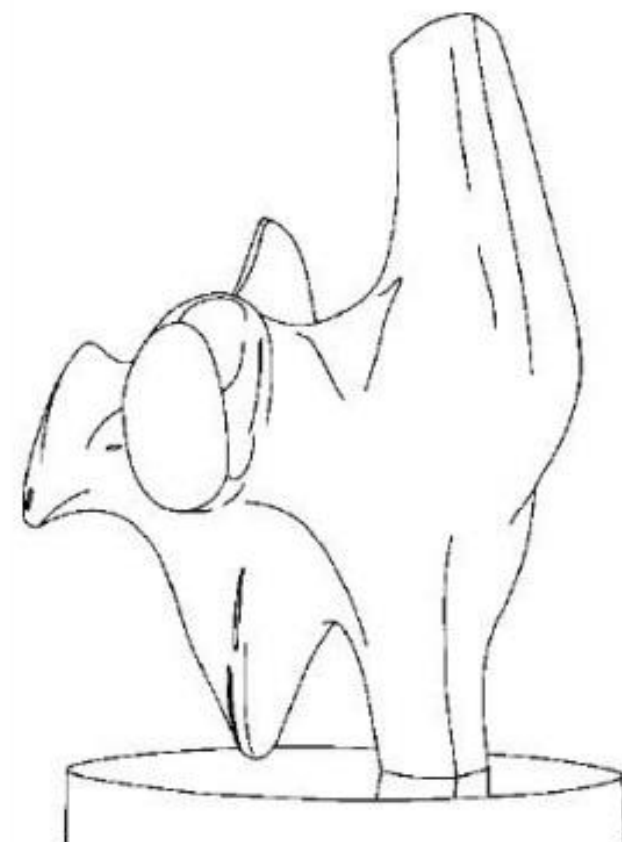


smoothed – original

Laplacian Pyramid

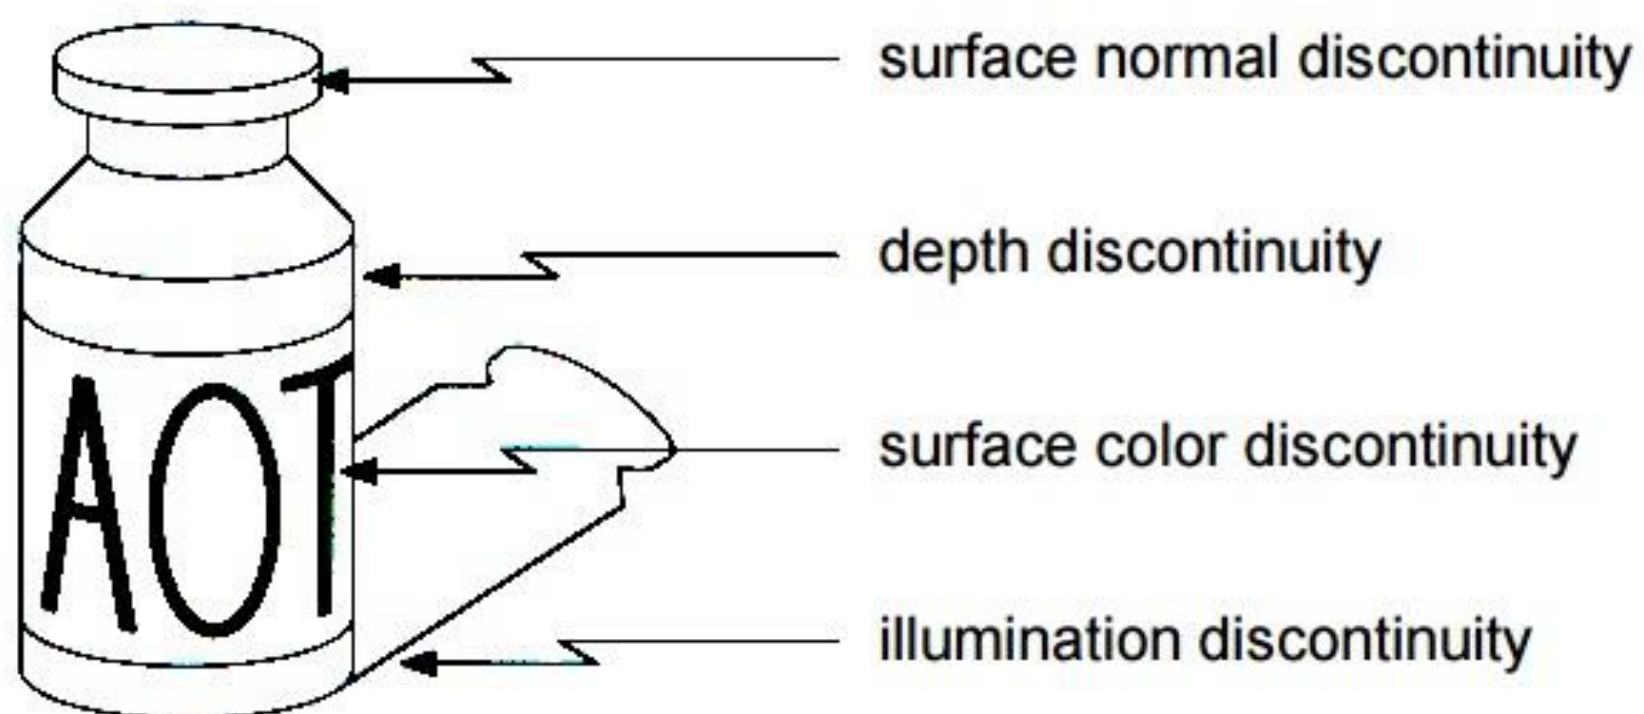


Edges in Depth



- Convert a 2D image into a set of curves
 - Extracts salient features of the scene
 - More compact than pixels

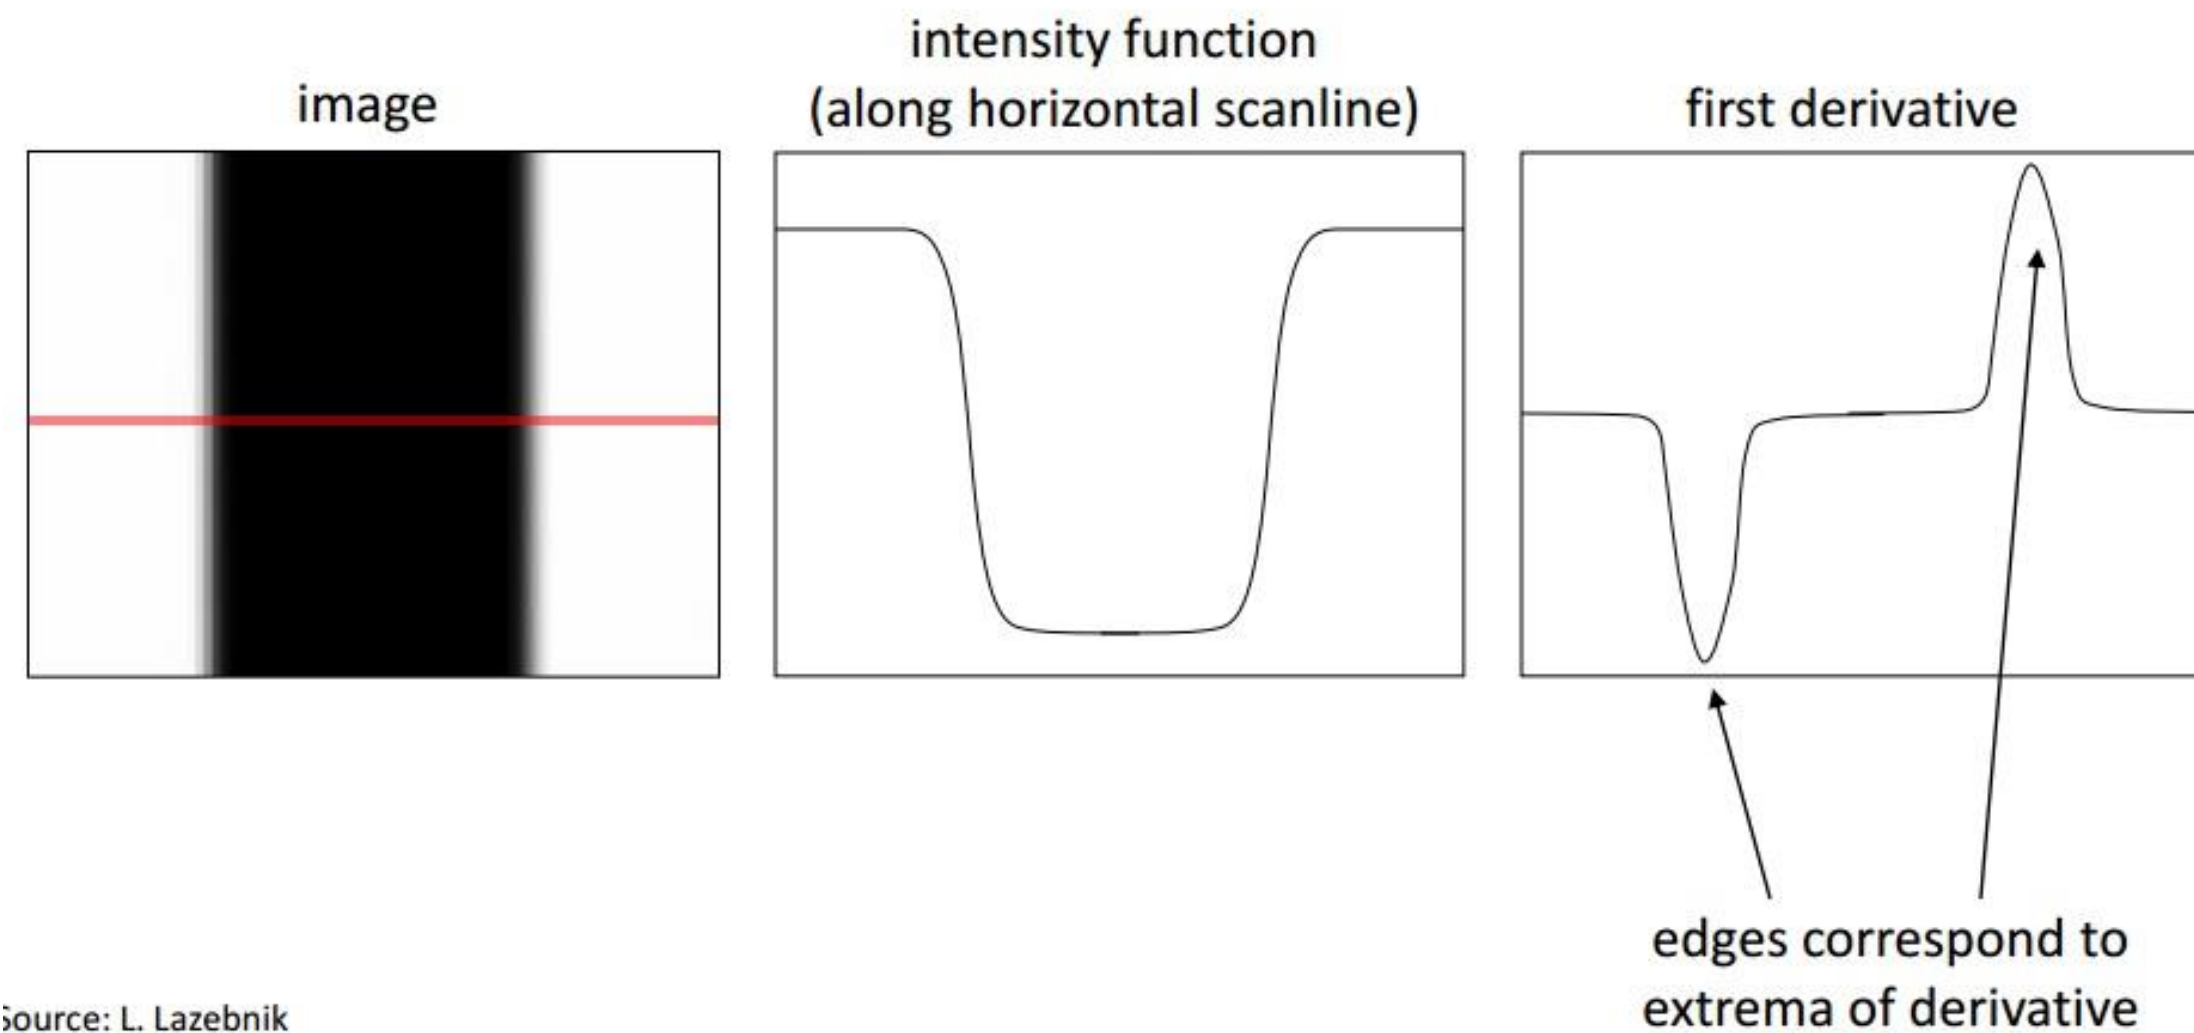
Edges in Depth



- Edges are caused by a variety of factors

Finding Edges

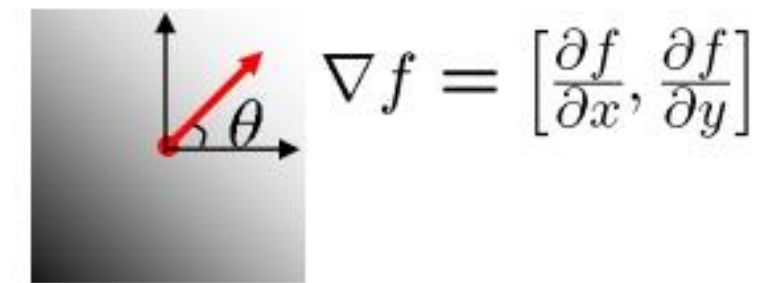
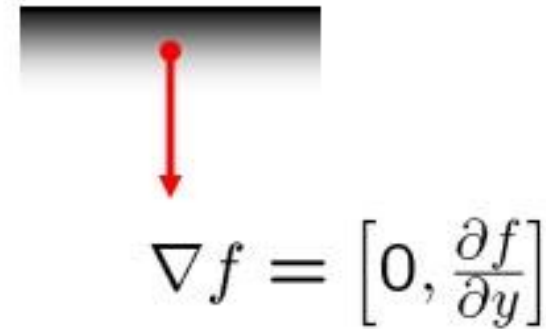
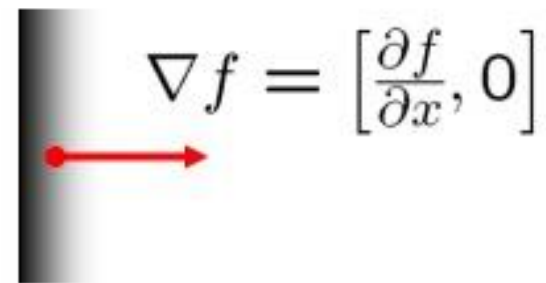
- An edge is a place of rapid change in the image intensity function



Edges as Gradient

- The *gradient* of an image: $\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$

The gradient points in the direction of most rapid increase in intensity



The *edge strength* is given by the gradient magnitude:

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$

The gradient direction is given by:

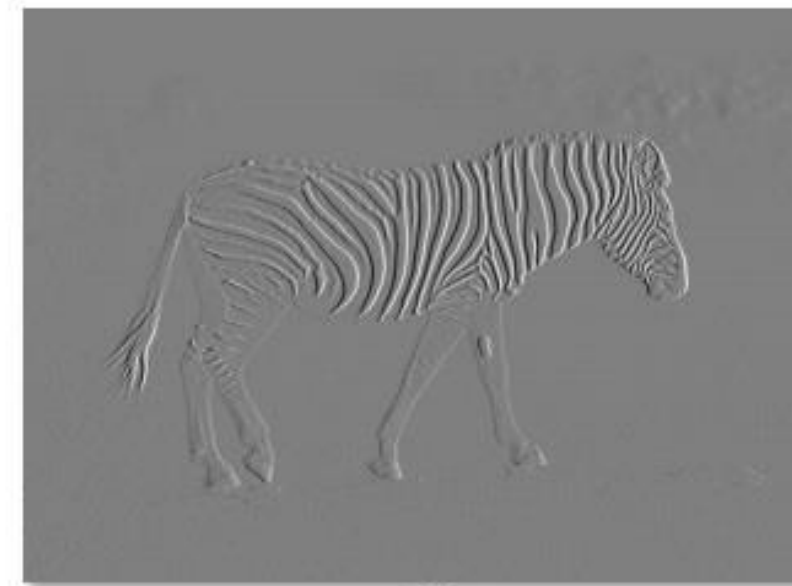
$$\theta = \tan^{-1} \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$

- how does this relate to the direction of the edge?

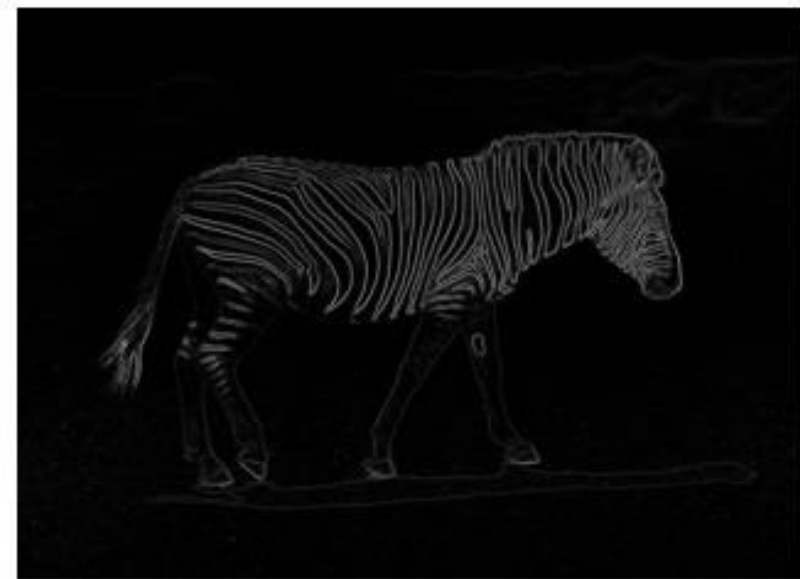
Edges as Gradient



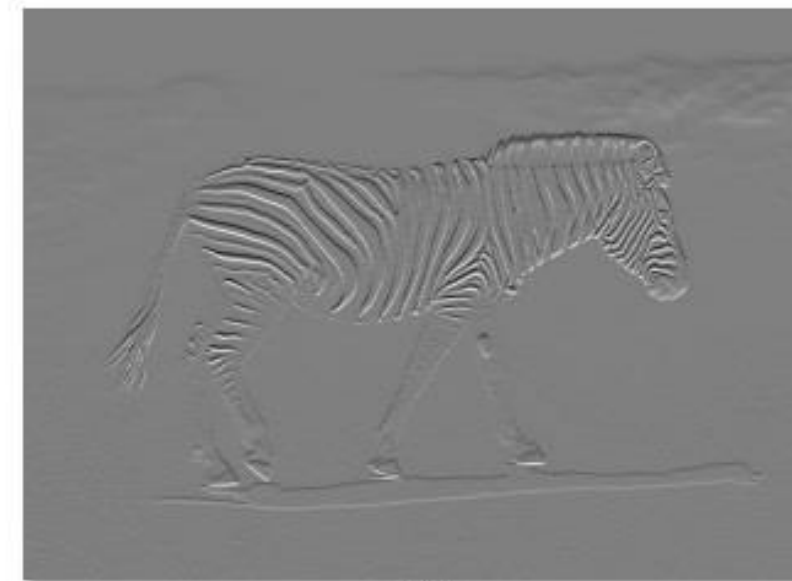
f



$\frac{\partial f}{\partial x}$

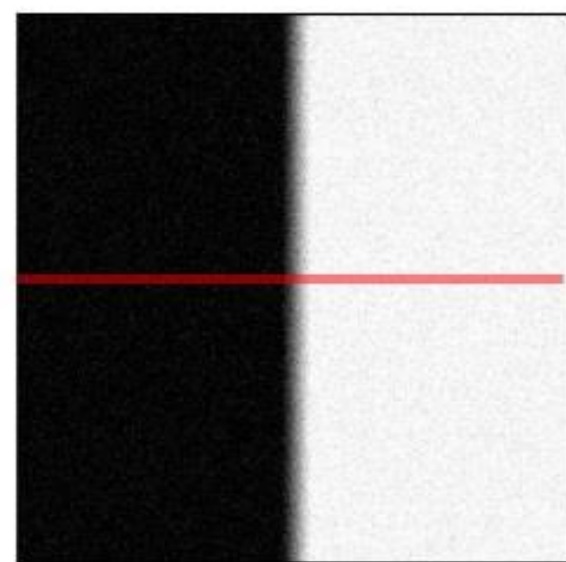


$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$

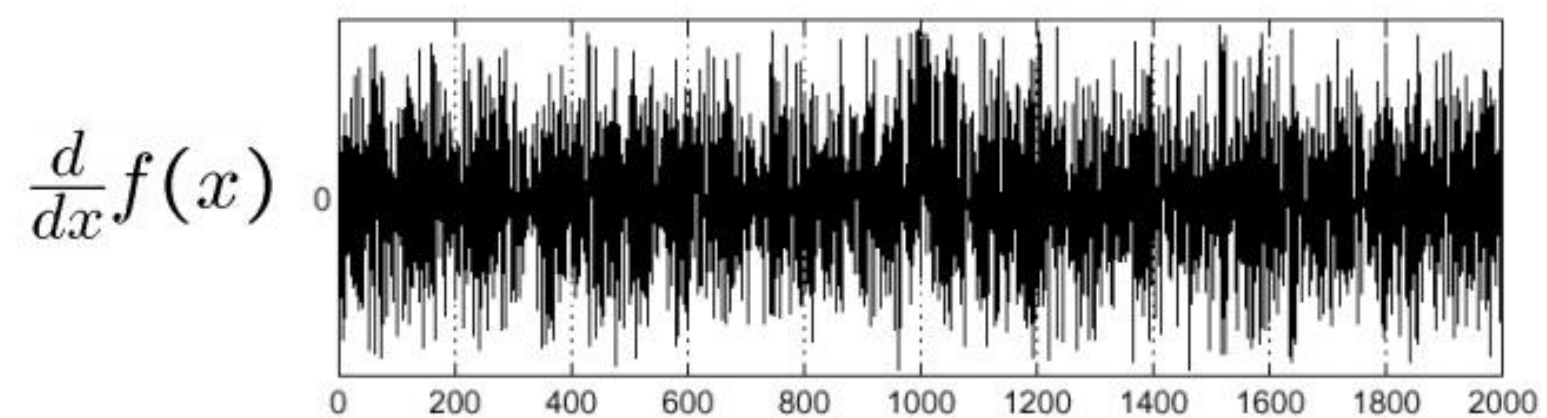
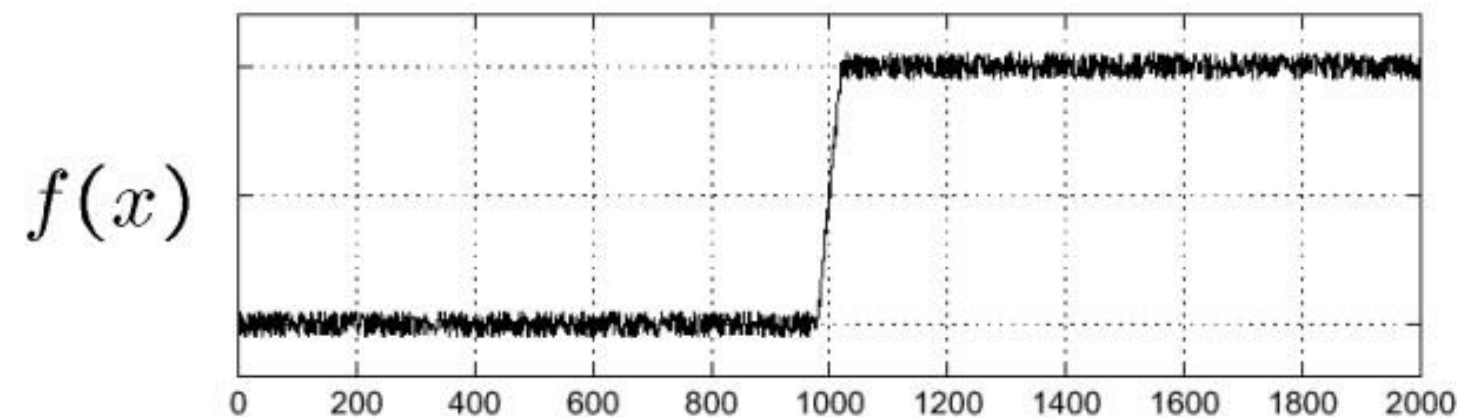


$\frac{\partial f}{\partial y}$

Edges for Noisy Image!



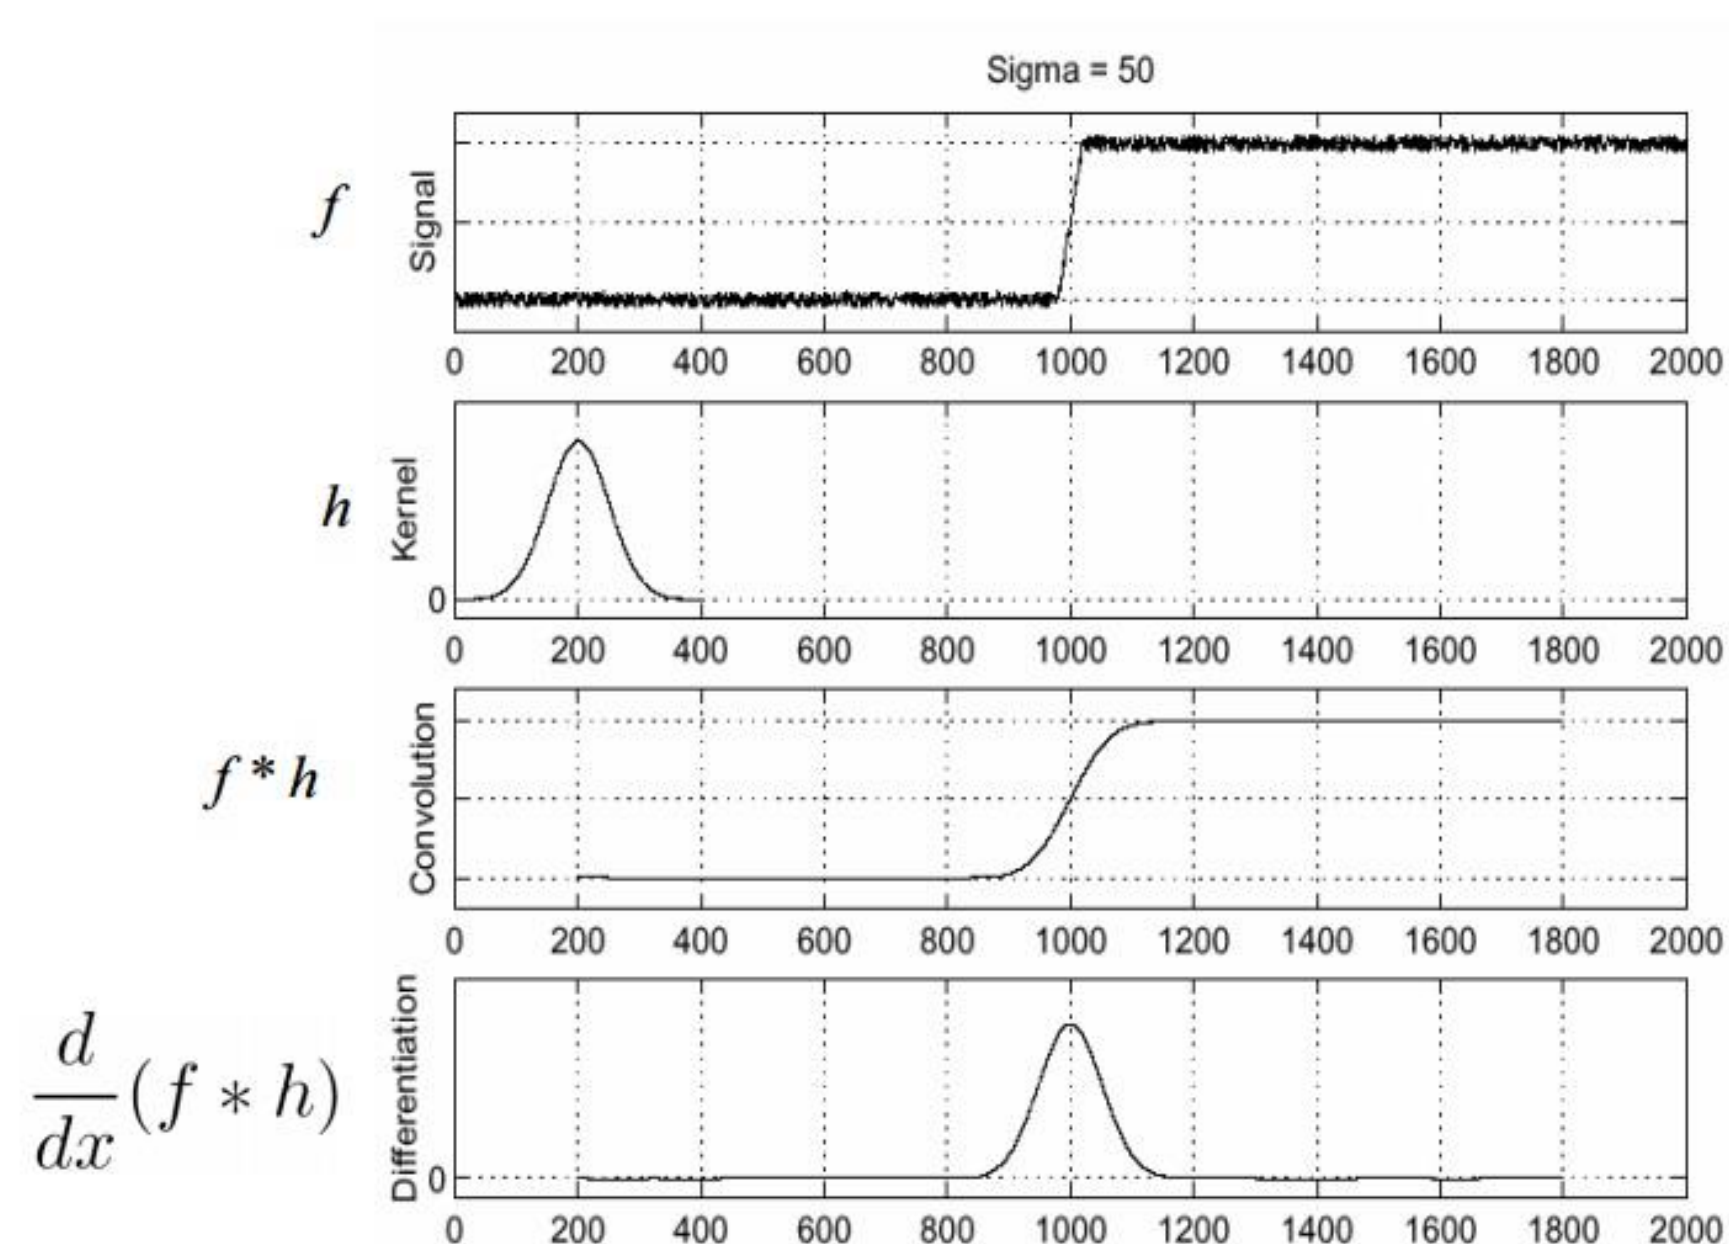
Noisy input image



Where is the edge?

Source: S. Seitz

Finding Edges in Noisy Images

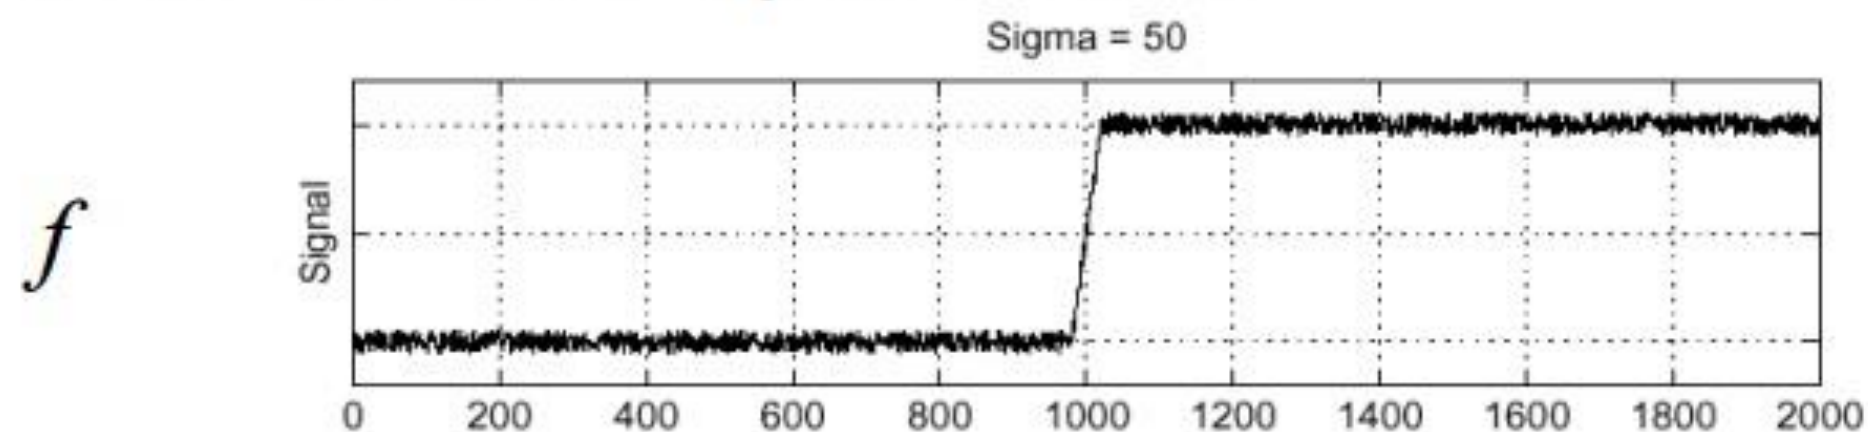


To find edges, look for peaks in $\frac{d}{dx}(f * h)$

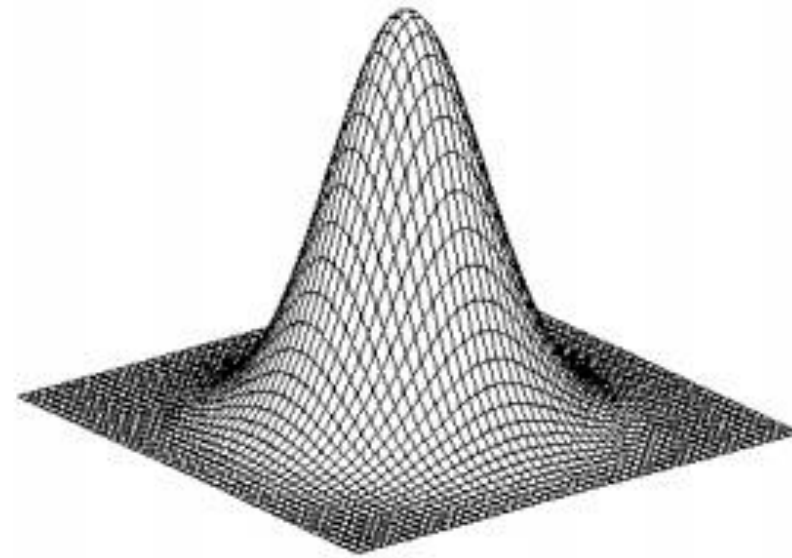
Differential and Associative Properties

$$\frac{d}{dx}(f * h) = f * \frac{d}{dx}h$$

- This saves us one operation:

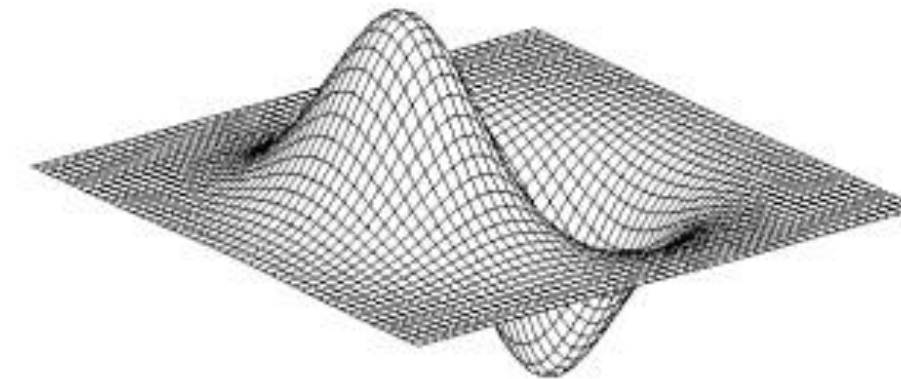


Derivative of Gaussian



Gaussian

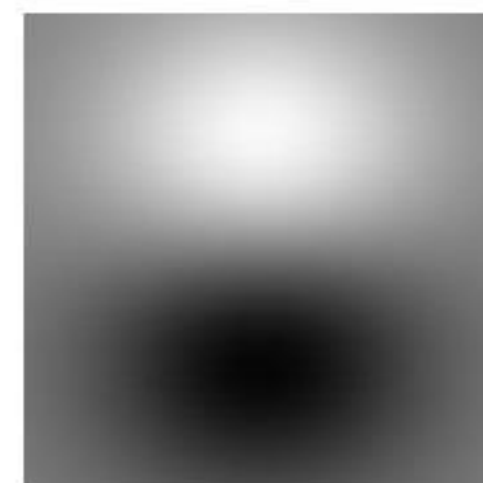
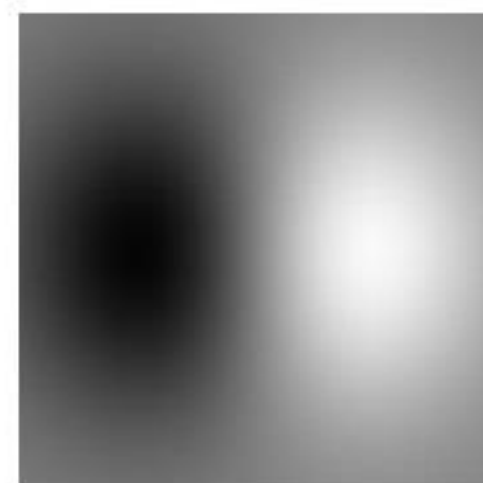
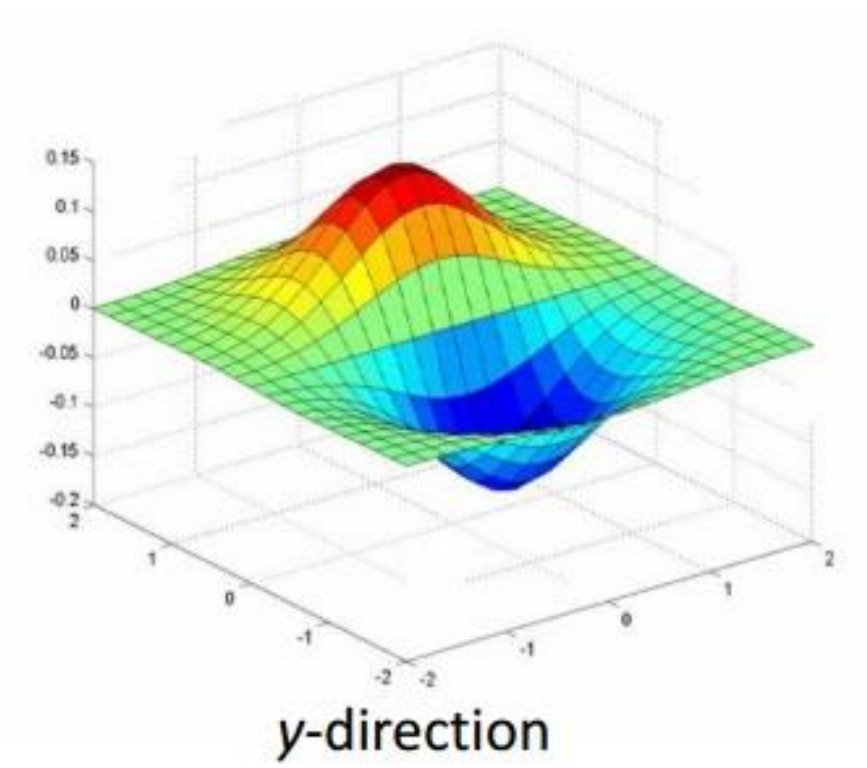
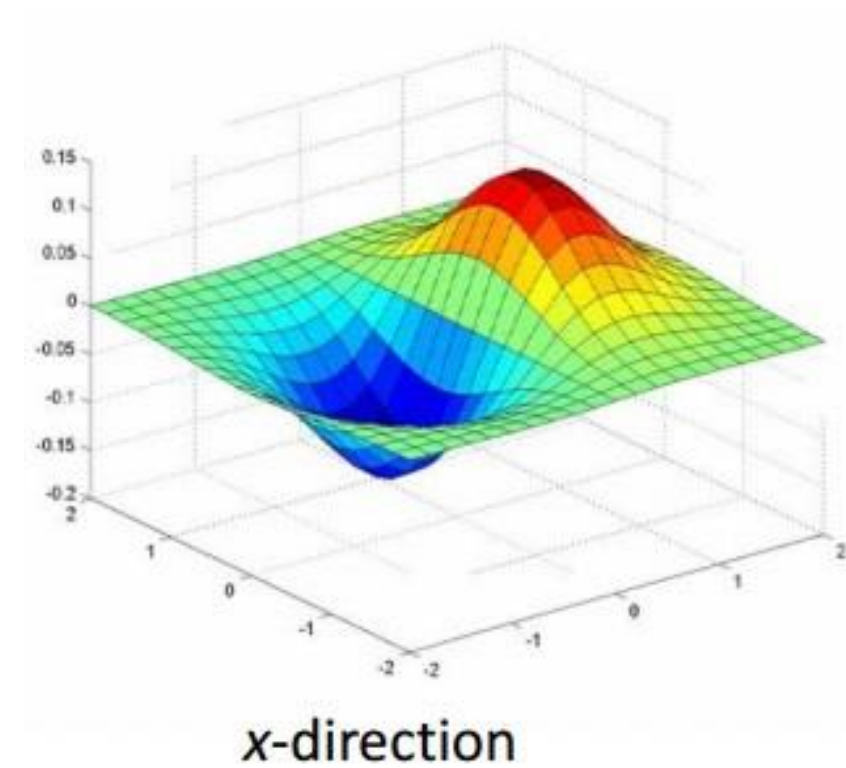
$$h_{\sigma}(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$



derivative of Gaussian (x)

$$\frac{\partial}{\partial x} h_{\sigma}(u, v)$$

Derivative of Gaussian



Sobel Filters

- Common approximation of derivative of Gaussian

-1	0	1
-2	0	2
-1	0	1

s_x

1	2	1
0	0	0
-1	-2	-1

s_y

Sobel Filter - Example



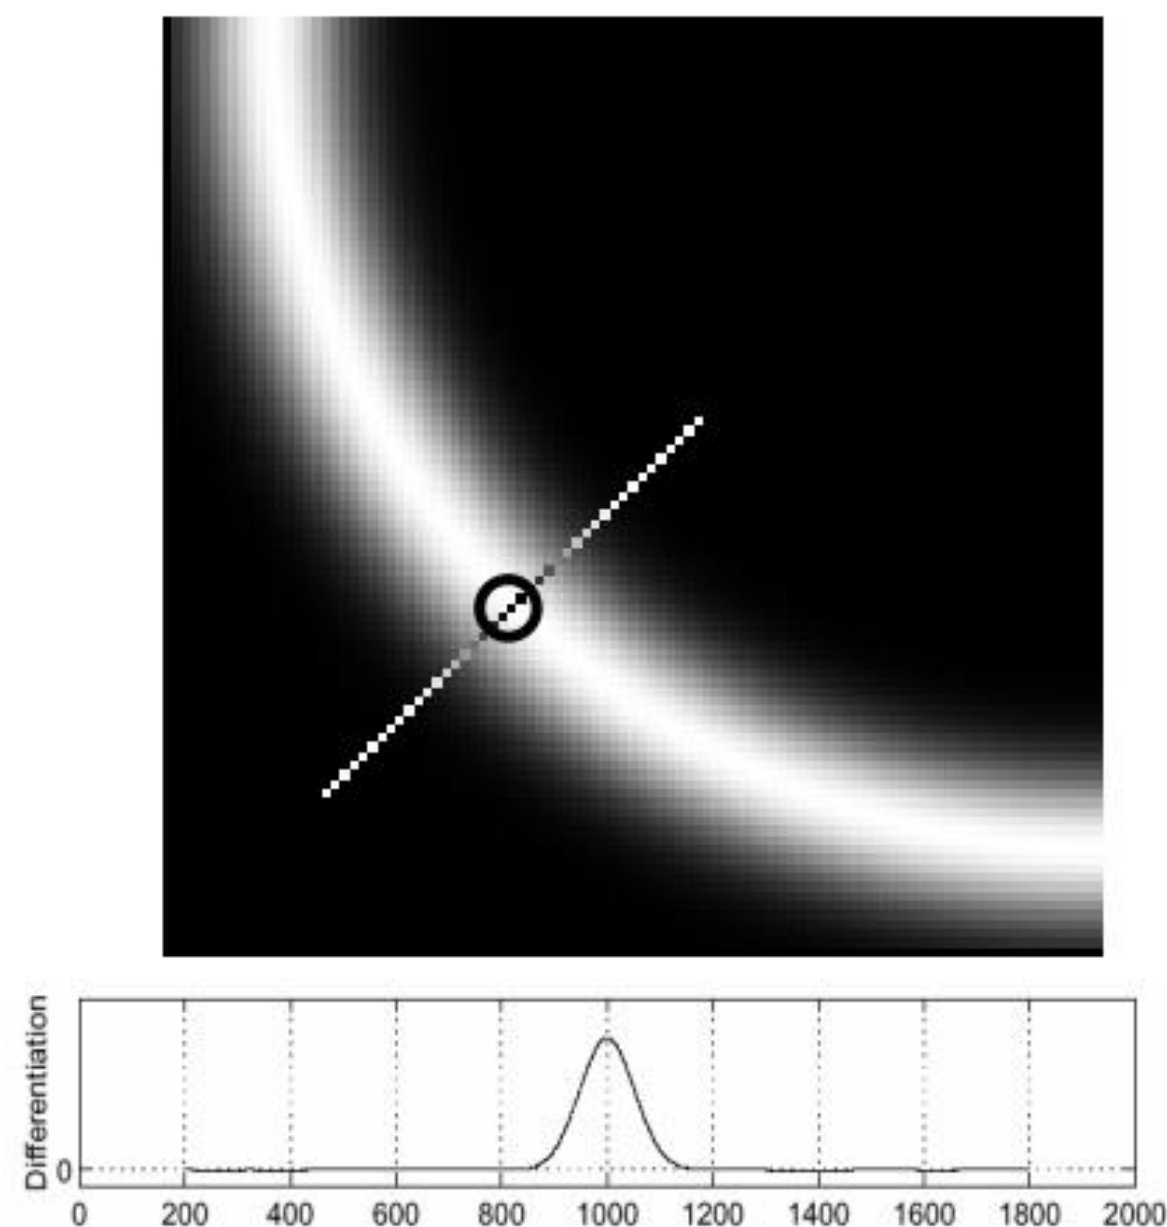
Source: Wikipedia

Sobel Filter - Example



thresholding

Non-Maxima Suppression



- Check if pixel is local maximum along gradient direction

Edge Thinning



thinning
(non-maximum suppression)